

# etcd: mission critical key-value store

Gopherfest  
16 May 2016

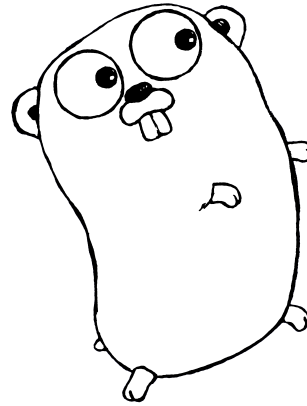
Gyu-Ho Lee  
CoreOS

# Welcome

Slides are here:

[github.com/gyuho/presentations](https://github.com/gyuho/presentations) (<https://github.com/gyuho/presentations>)

# Go + etcd



# Agenda

- What is etcd
- Why
- Go
- Q/A

# What is etcd

## etcd is ...

- Distributed key-value store
- Open source [github.com/coreos/etcd](https://github.com/coreos/etcd) (<https://github.com/coreos/etcd>) (~ June 2013)
- Still new, compared to ZooKeeper (~ May 2008)
- Google [Kubernetes](http://kubernetes.io/), YouTube [Doorman](https://github.com/youtube/doorman), ...
- Red Hat, EMC, Cisco, Huawei, Baidu, Alibaba...

# etcd API

```
cli.Put(ctx, "foo", "bar", Lease)
cli.Get(ctx, "foo")
cli.Delete(ctx, "foo")

// Transaction
kvc.Txn(ctx).
  If(clientv3.Compare(clientv3.Value("key"), ">", "abc")). // txn value comparisons are lexical
  Then(clientv3.OpPut("key", "XYZ")).                      // this runs, since 'xyz' > 'abc'
  Else(clientv3.OpPut("key", "ABC")).
  Commit()

// Watch for updates on key
ch := cli.Watch(ctx, "foo")
for res := range ch {}

// Distributed locks
mu := concurrency.NewMutex(cli, "foo")
mu.Lock()
mu.Unlock()
```

# Why etcd



# Use etcd to store configuration

For small chunks of data

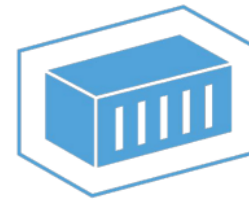
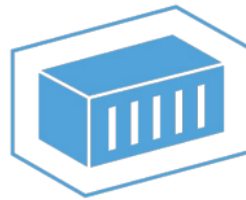
```
maxReqBytes = 1.5 * 1024 * 1024 // 1.5MB
```

```
DefaultQuotaBytes = int64(2 * 1024 * 1024 * 1024) // 2GB
```

```
MaxQuotaBytes = int64(8 * 1024 * 1024 * 1024) // 8GB
```

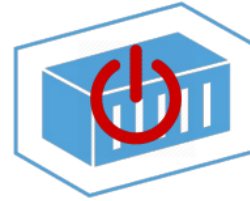
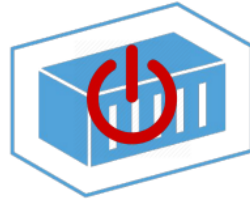
# Updates

Security updates?



How would you update the cluster of machines?

# Traditional way



- Reboot with downtime
- Too Manual

# CoreOS updates with etcd

If you run your application on CoreOS,  
your OS gets Automatic, No-downtime updates

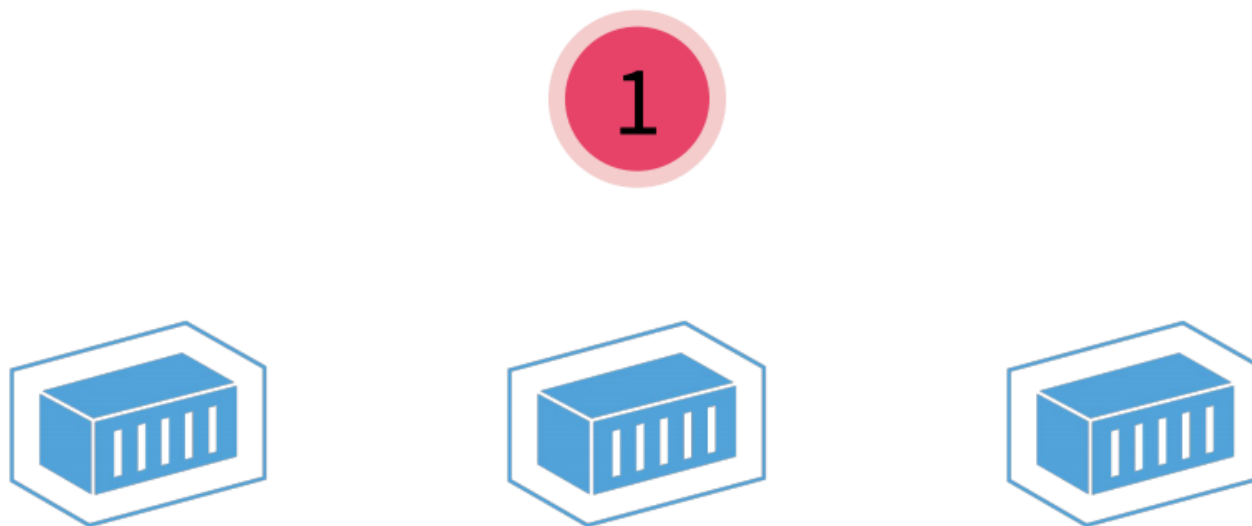
**1032.1.0** Release Date: May 5, 2016   kernel: 4.5.2   rkt: 1.2.1   docker: 1.10.3   etcd: 0.4.9, 2.3.2   fleet: 0.11.7   systemd: 229

## Security Updates:

- OpenSSL [1.0.2h](#) for [CVE-2016-2105](#), [CVE-2016-2106](#), [CVE-2016-2107](#), [CVE-2016-2109](#), [CVE-2016-2176](#)
- ntpd [4.2.8p7](#) for [CVE-2016-1551](#), [CVE-2016-1549](#), [CVE-2016-2516](#), [CVE-2016-2517](#), [CVE-2016-2518](#), [CVE-2016-2519](#), [CVE-2016-1547](#), [CVE-2016-1548](#), [CVE-2015-7704](#), [CVE-2015-8138](#), [CVE-2016-1550](#)
- git [2.7.3-r1](#) for [CVE-2015-7545](#), [CVE-2016-2315](#), [CVE-2016-2315](#)
- jq [1.5-r2](#) for [CVE-2015-8863](#)

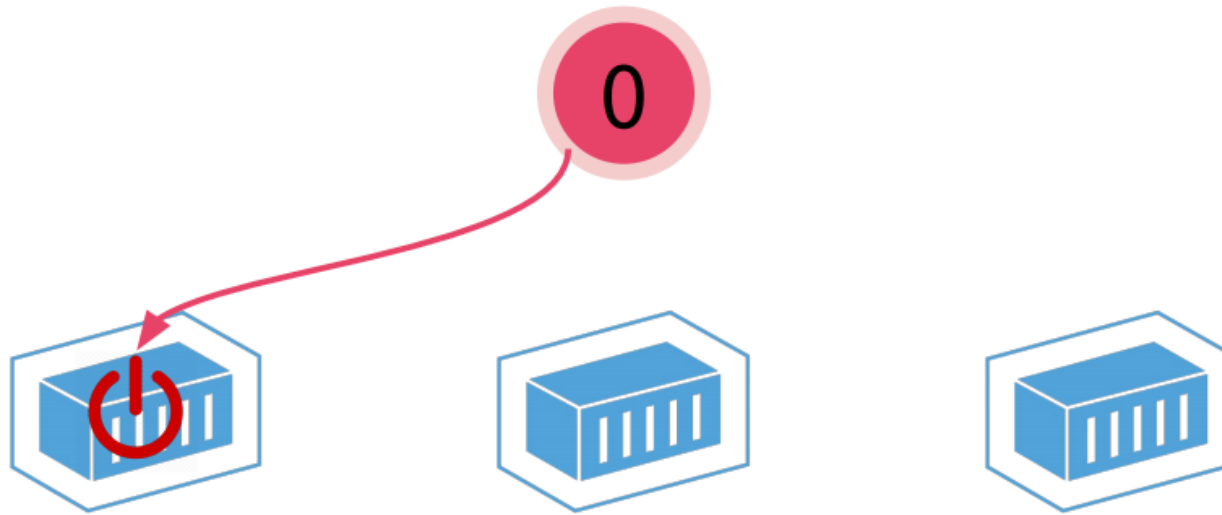
# CoreOS updates with etcd

- CoreOS updates are done by [locksmith](https://github.com/coreos/locksmith) (<https://github.com/coreos/locksmith>)
- locksmith is built on top of etcd
- locksmith stores semaphore values in etcd
- ensure that only subset of cluster are rebooting at any given time



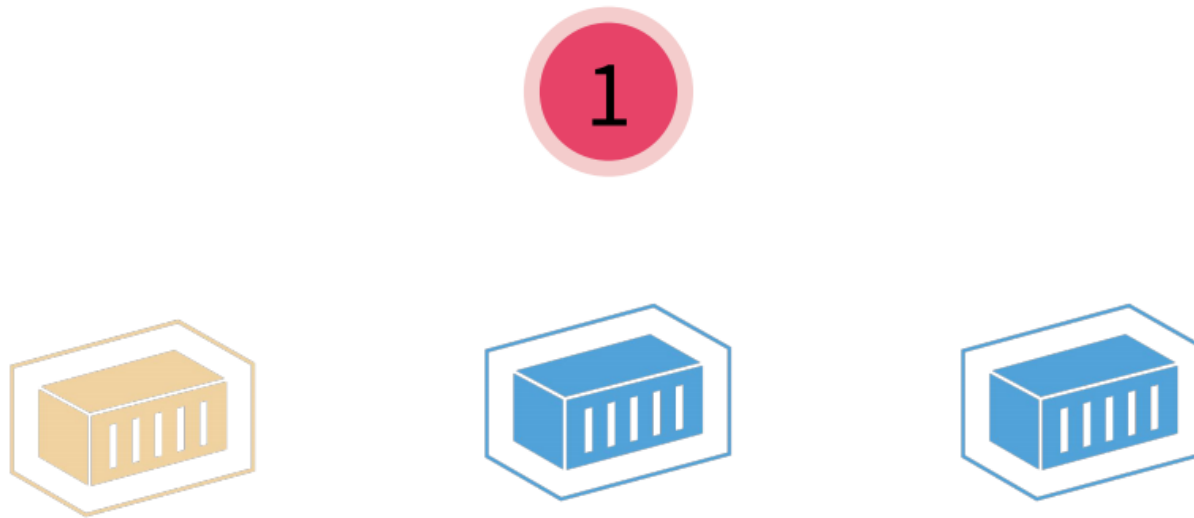
# CoreOS updates with etcd

Decrement semaphore when rebooting



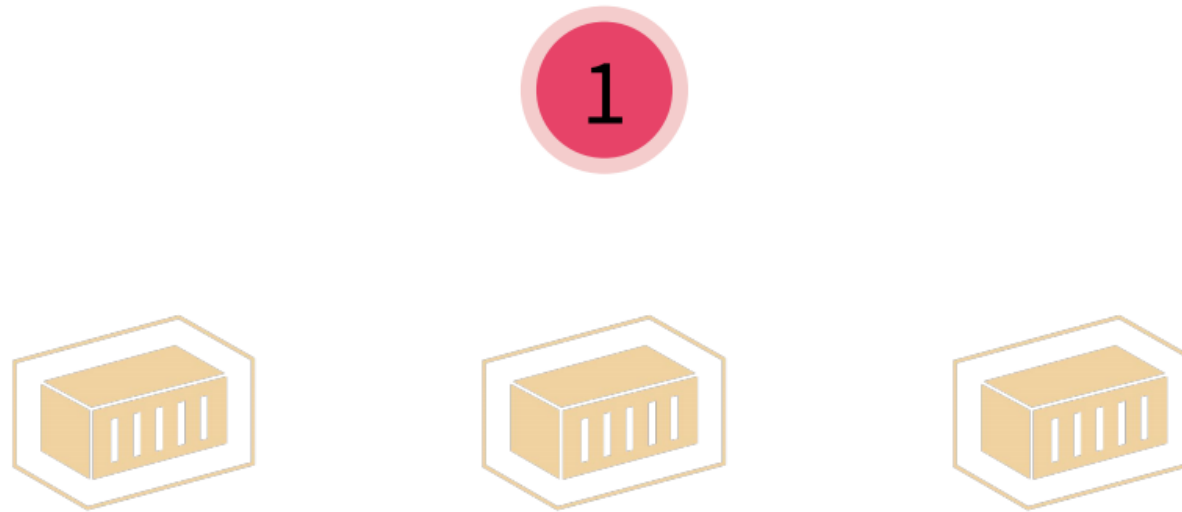
# CoreOS updates with etcd

- Automatic
- No downtime



# CoreOS updates with etcd

Your cluster is now secured

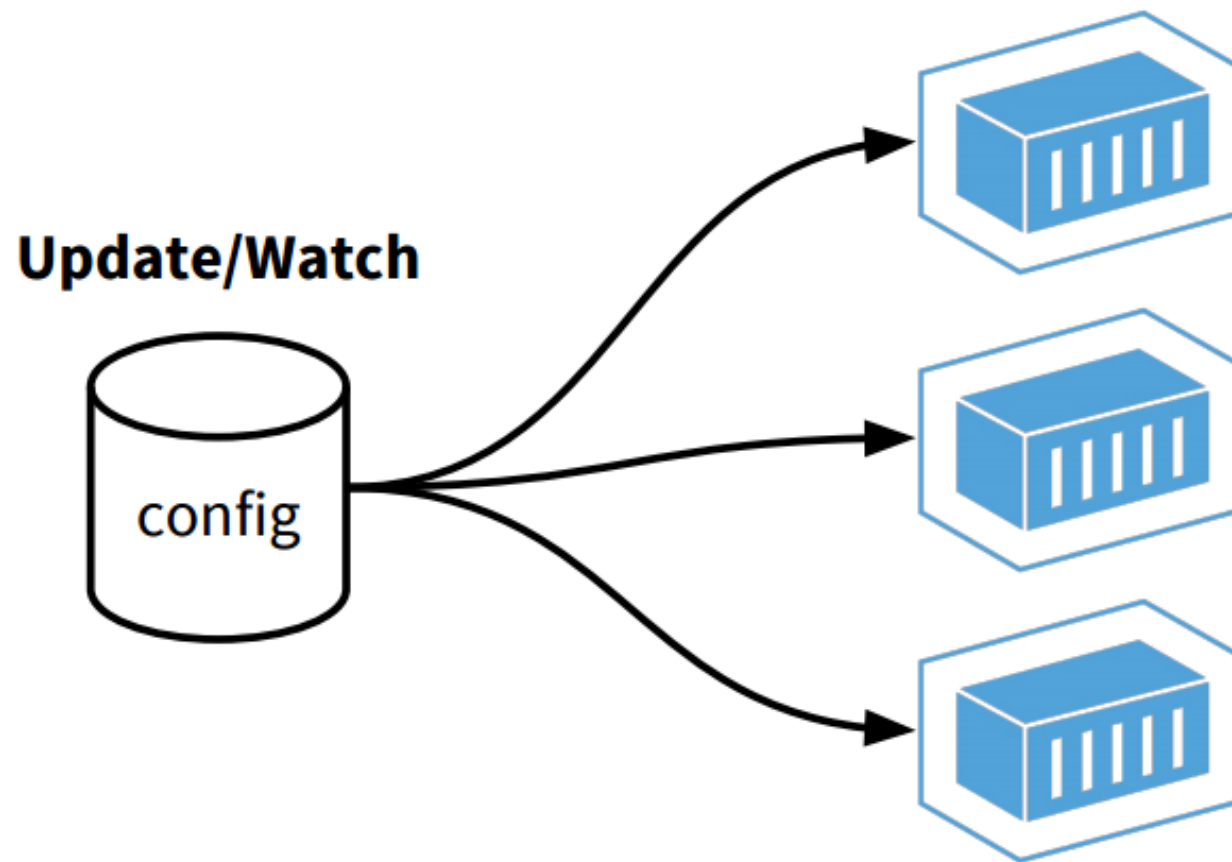




# Use etcd for "critical" configuration

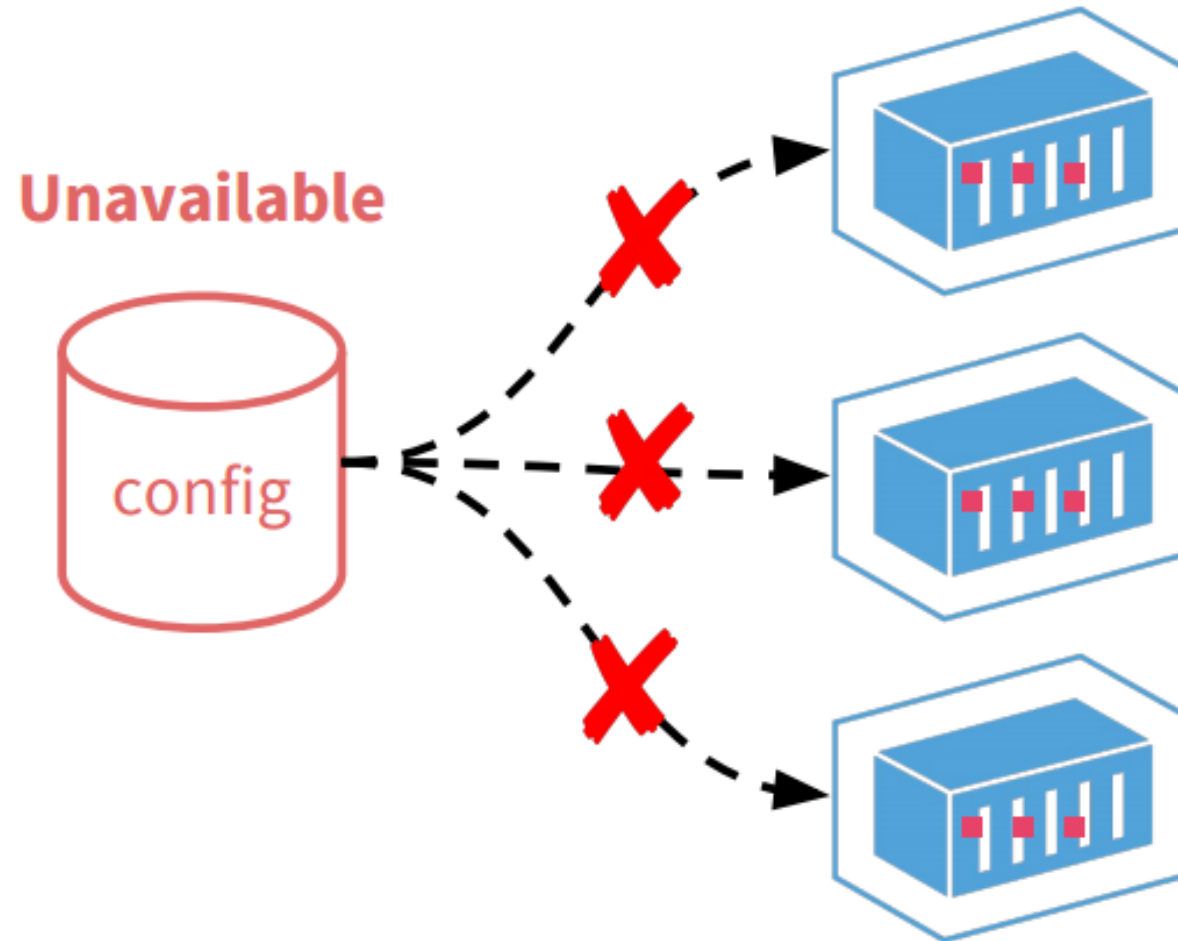
## Bad practice

What if this machine goes down?



# Bad practice


Single point of failure



## Good practice

etcd replicate your data to multiple machine

and still provides consistent view of your data

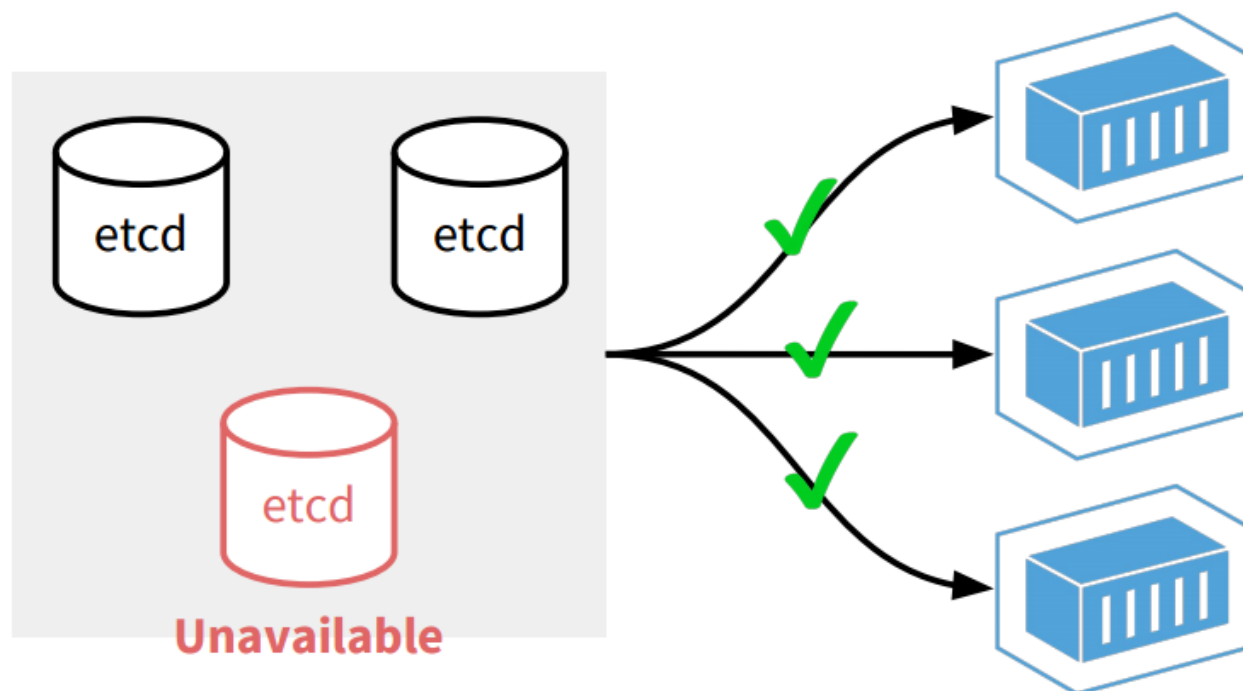


## Good practice

etcd can tolerate machine failures

Can tolerate up to 1 out of 3-node cluster

Can tolerate up to 2 out of 5-node cluster



# Demo

[play.etcd.io](http://play.etcd.io) (<http://play.etcd.io>)

Join me!

# etcd

## Consistent view of critical configuration

- Strong consistency (no stale reads)
- Different than eventual consistency (conflicts, latest timestamps wins)

## Highly available configuration store

- Resilient to a single point of failures & network partitions

## Watchable

- Push configuration updates to application

## Why not ZooKeeper or Consul?

They are all great projects.

They have their own use cases.

etcd is built for scalability and reliability.



# etcd Project Status: Performance

# etcd v3

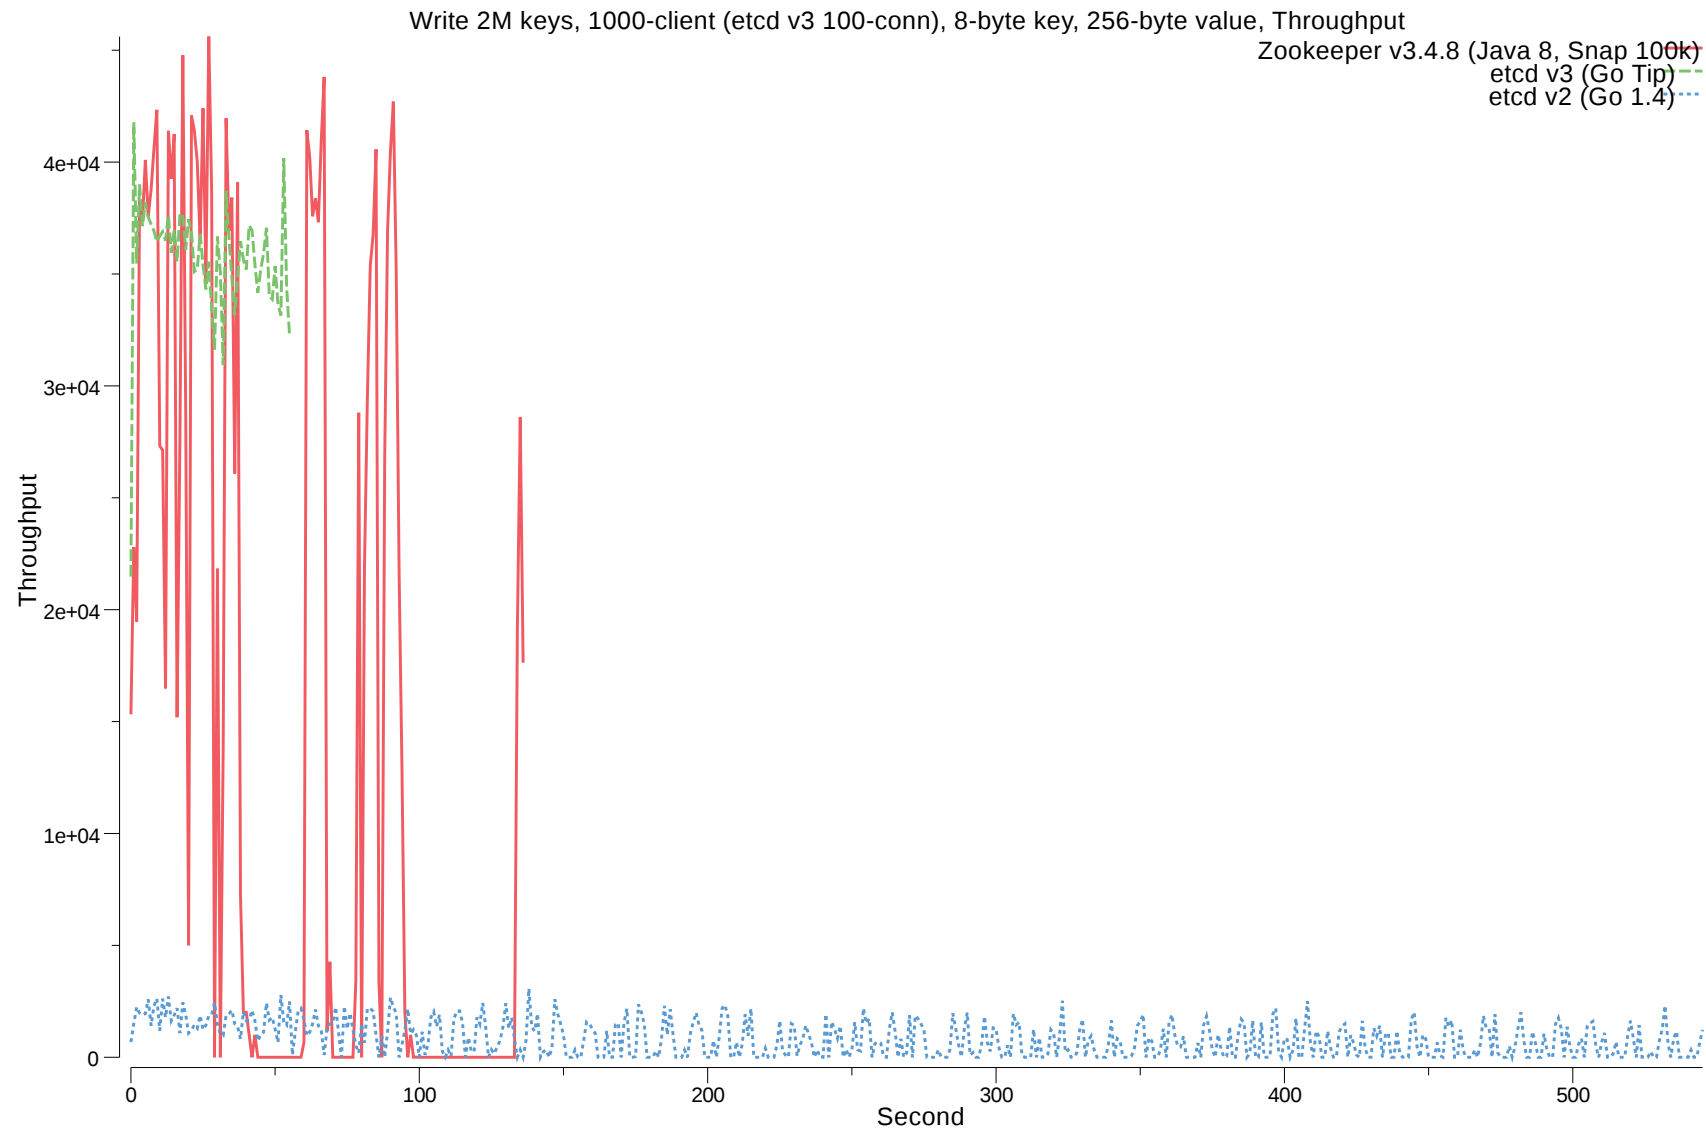
**BoltDB** (<https://github.com/boltdb/bolt>)

- B+tree disk storage
- Incremental snapshot
- vs. ZooKeeper snapCount 10,000

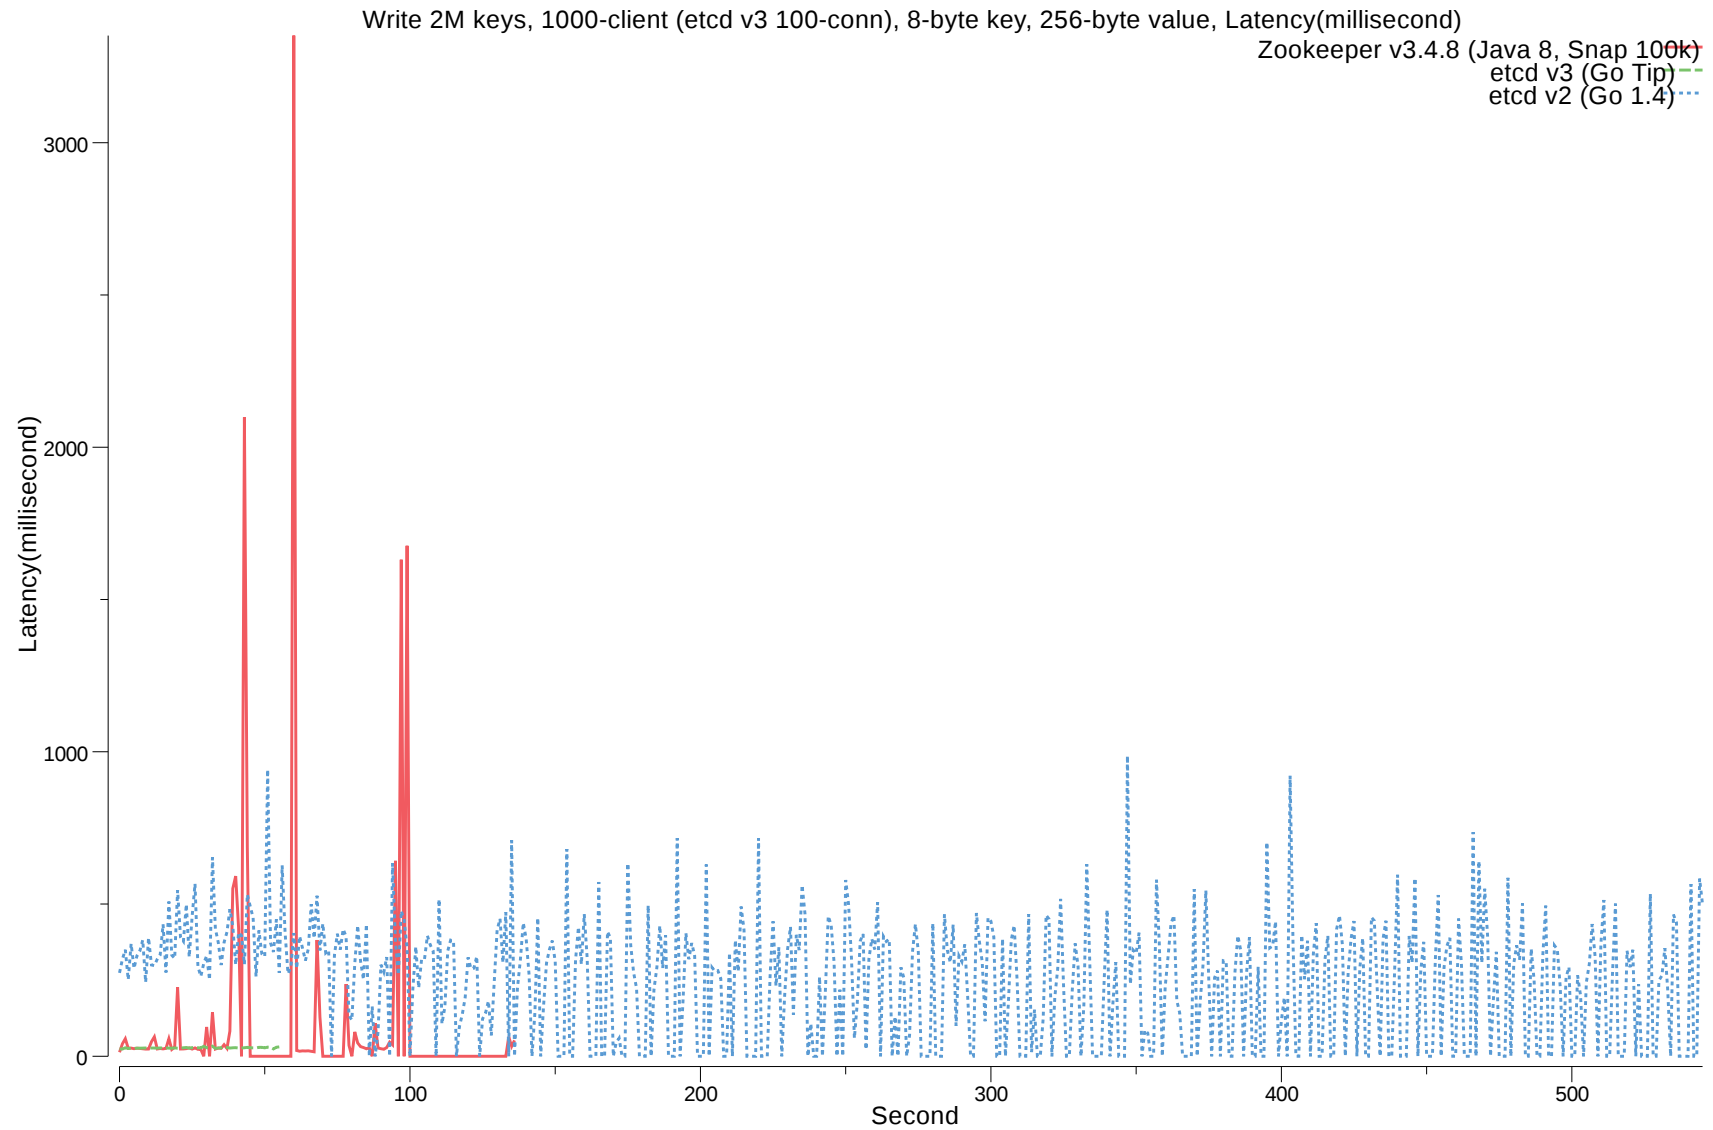
**gRPC** (<http://www.grpc.io/>)

- Protocol Buffer
- HTTP/2
- streams, less TCP congestions

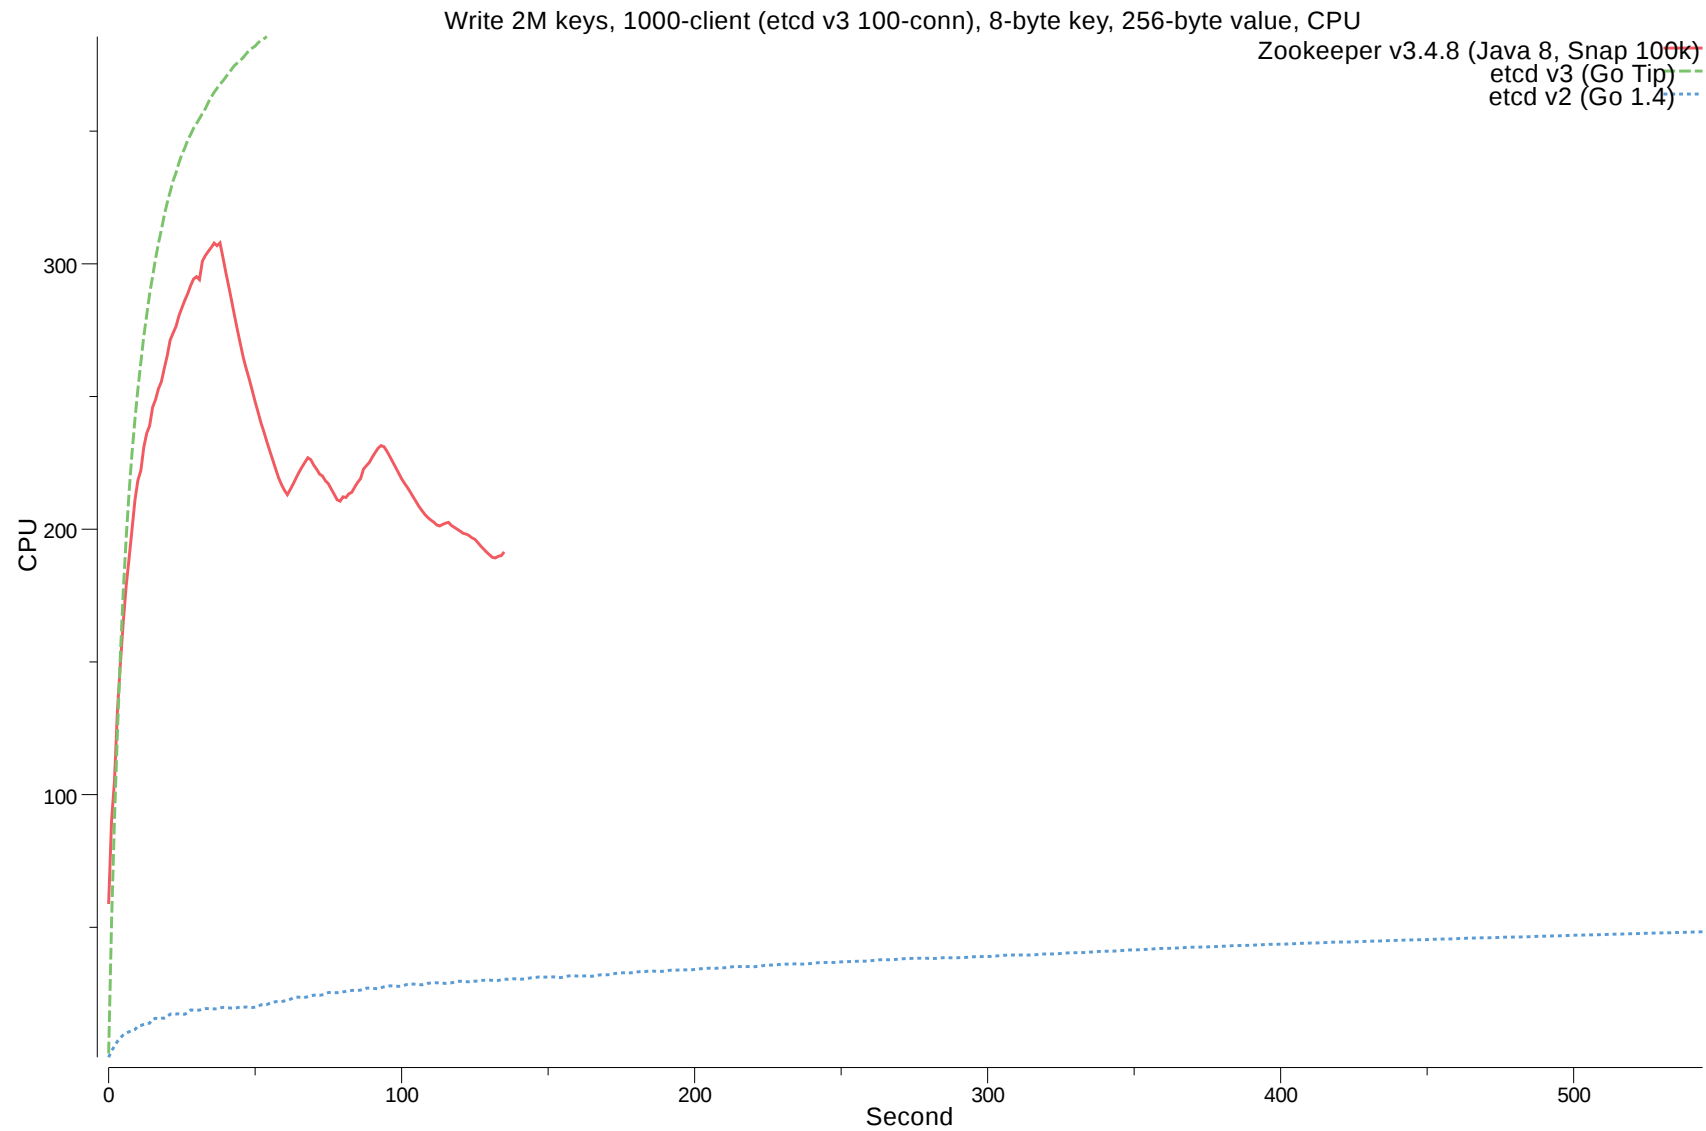
# Throughput



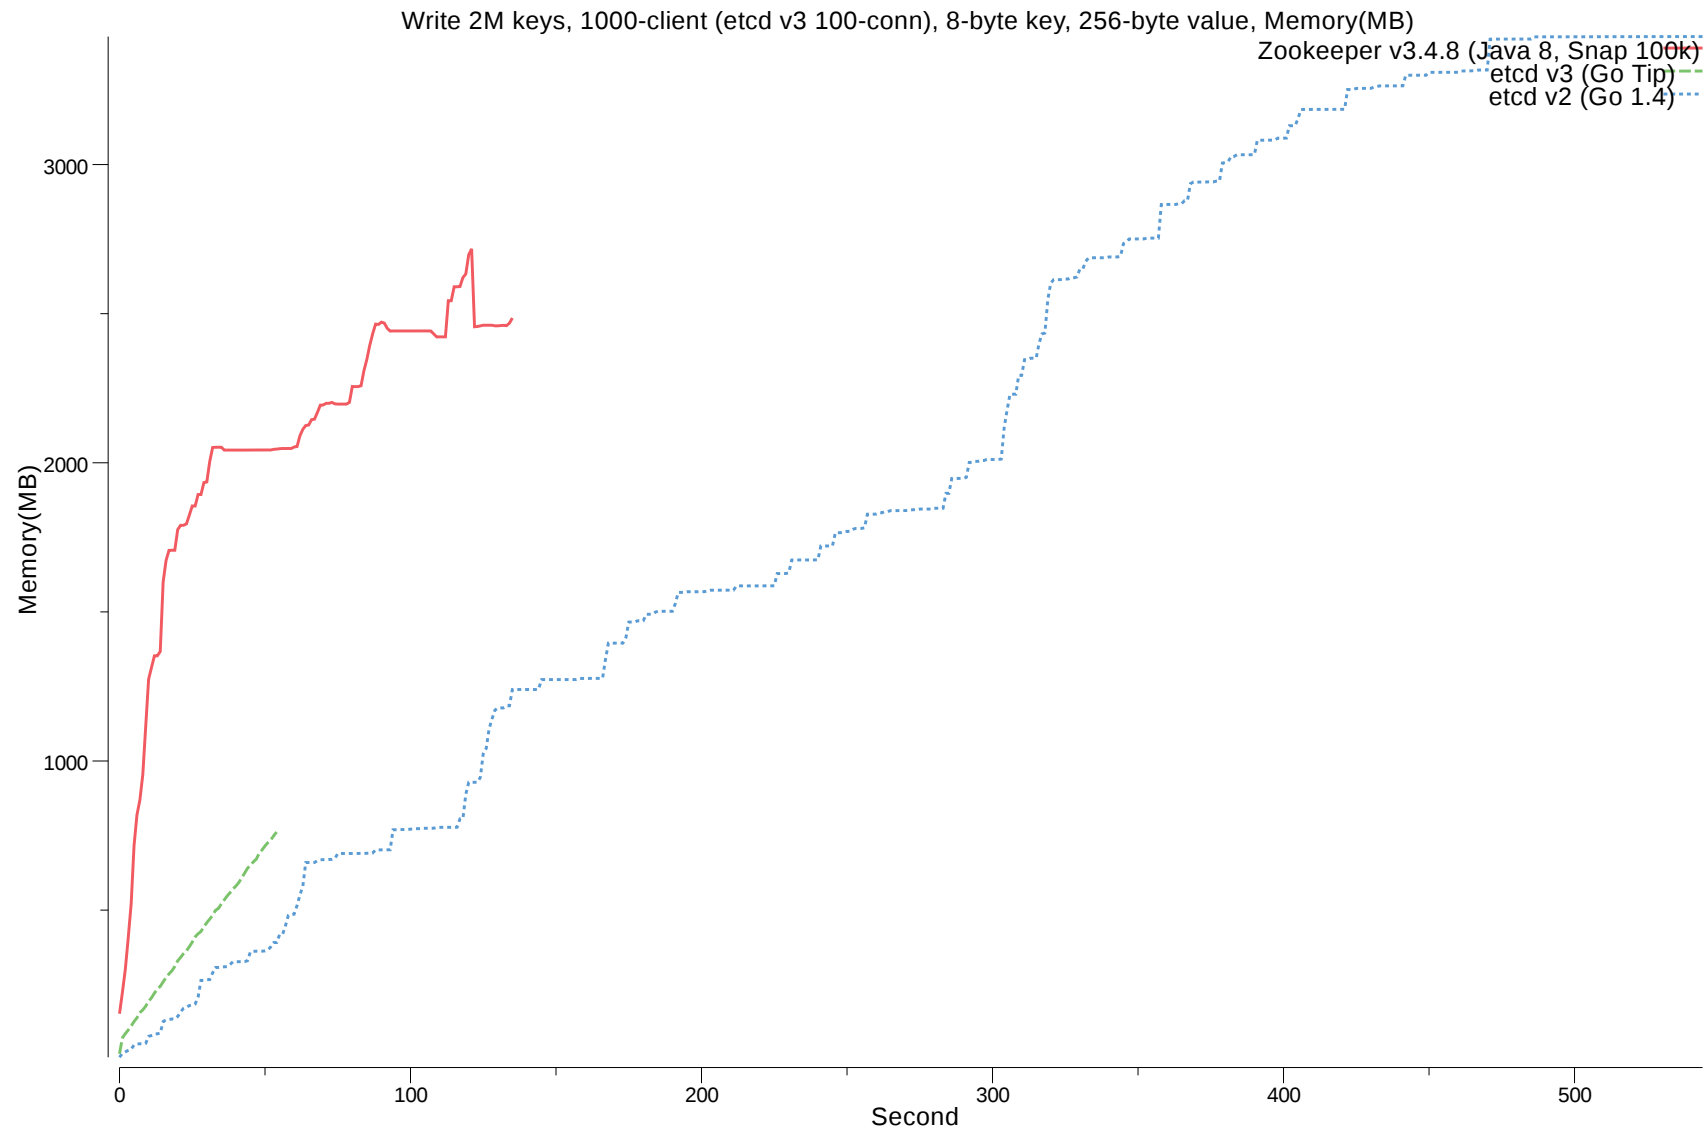
# Latency



# CPU

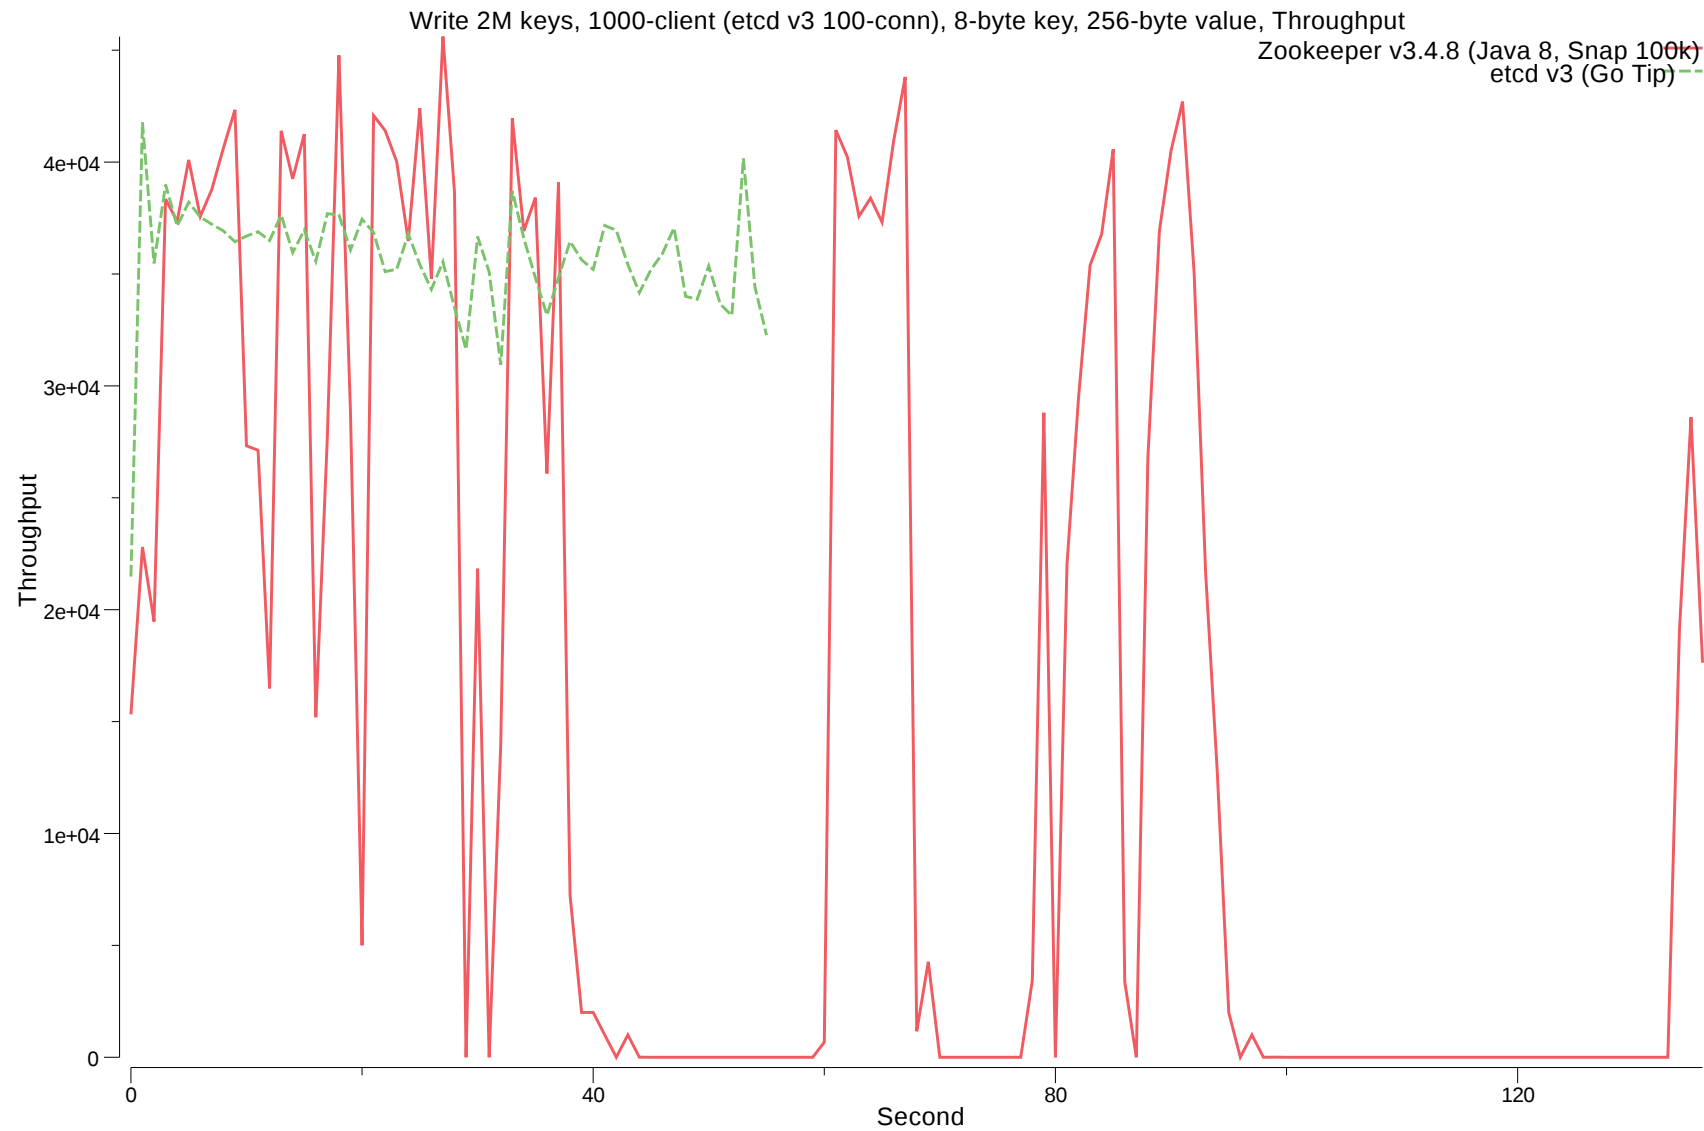


# Memory



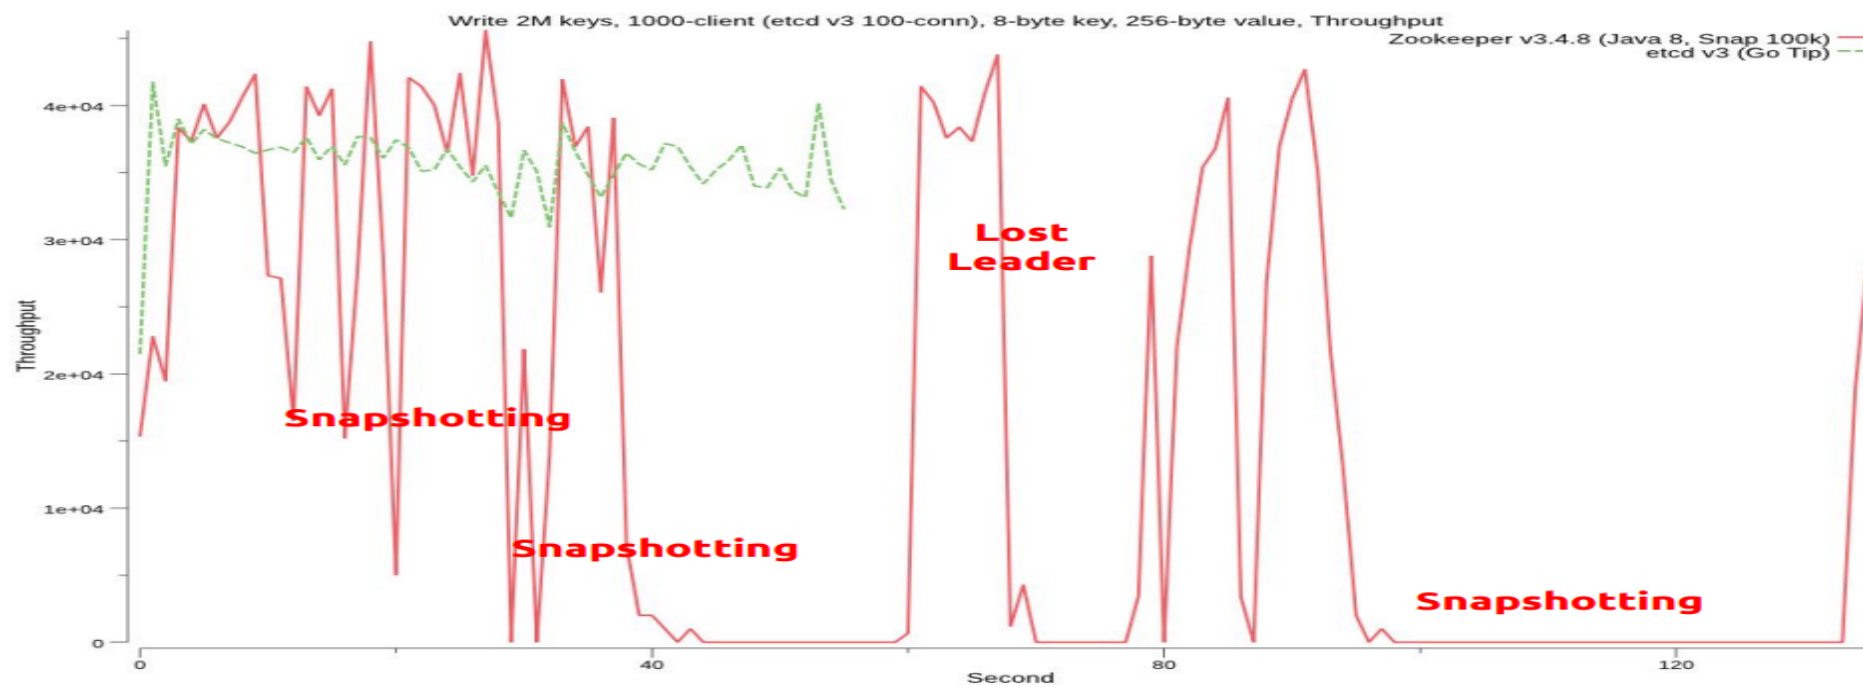
# etcd Project Status: Reliability

# Throughput





# etcd Reliability



## ZooKeeper logs

```
07:16:35 [Snapshot Thread:FileTxnSnapLog@240] - Snapshotting...
07:16:43 fsync-ing the write ahead log in SyncThread:3 took 1224ms...
07:16:46 fsync-ing the write ahead log in SyncThread:3 took 3205ms... // Snapshotting
...
07:17:14 [FastLeaderElection@818] - New election... // Leader Election
```

# etcd Reliability

Functional tests [dash.etcd.io](http://dash.etcd.io) (<http://dash.etcd.io/dashboard/db/functional-tests>)

- Kill one/all members
- Kill leader
- Network partition
- Network latency

Consistency checking after recovery

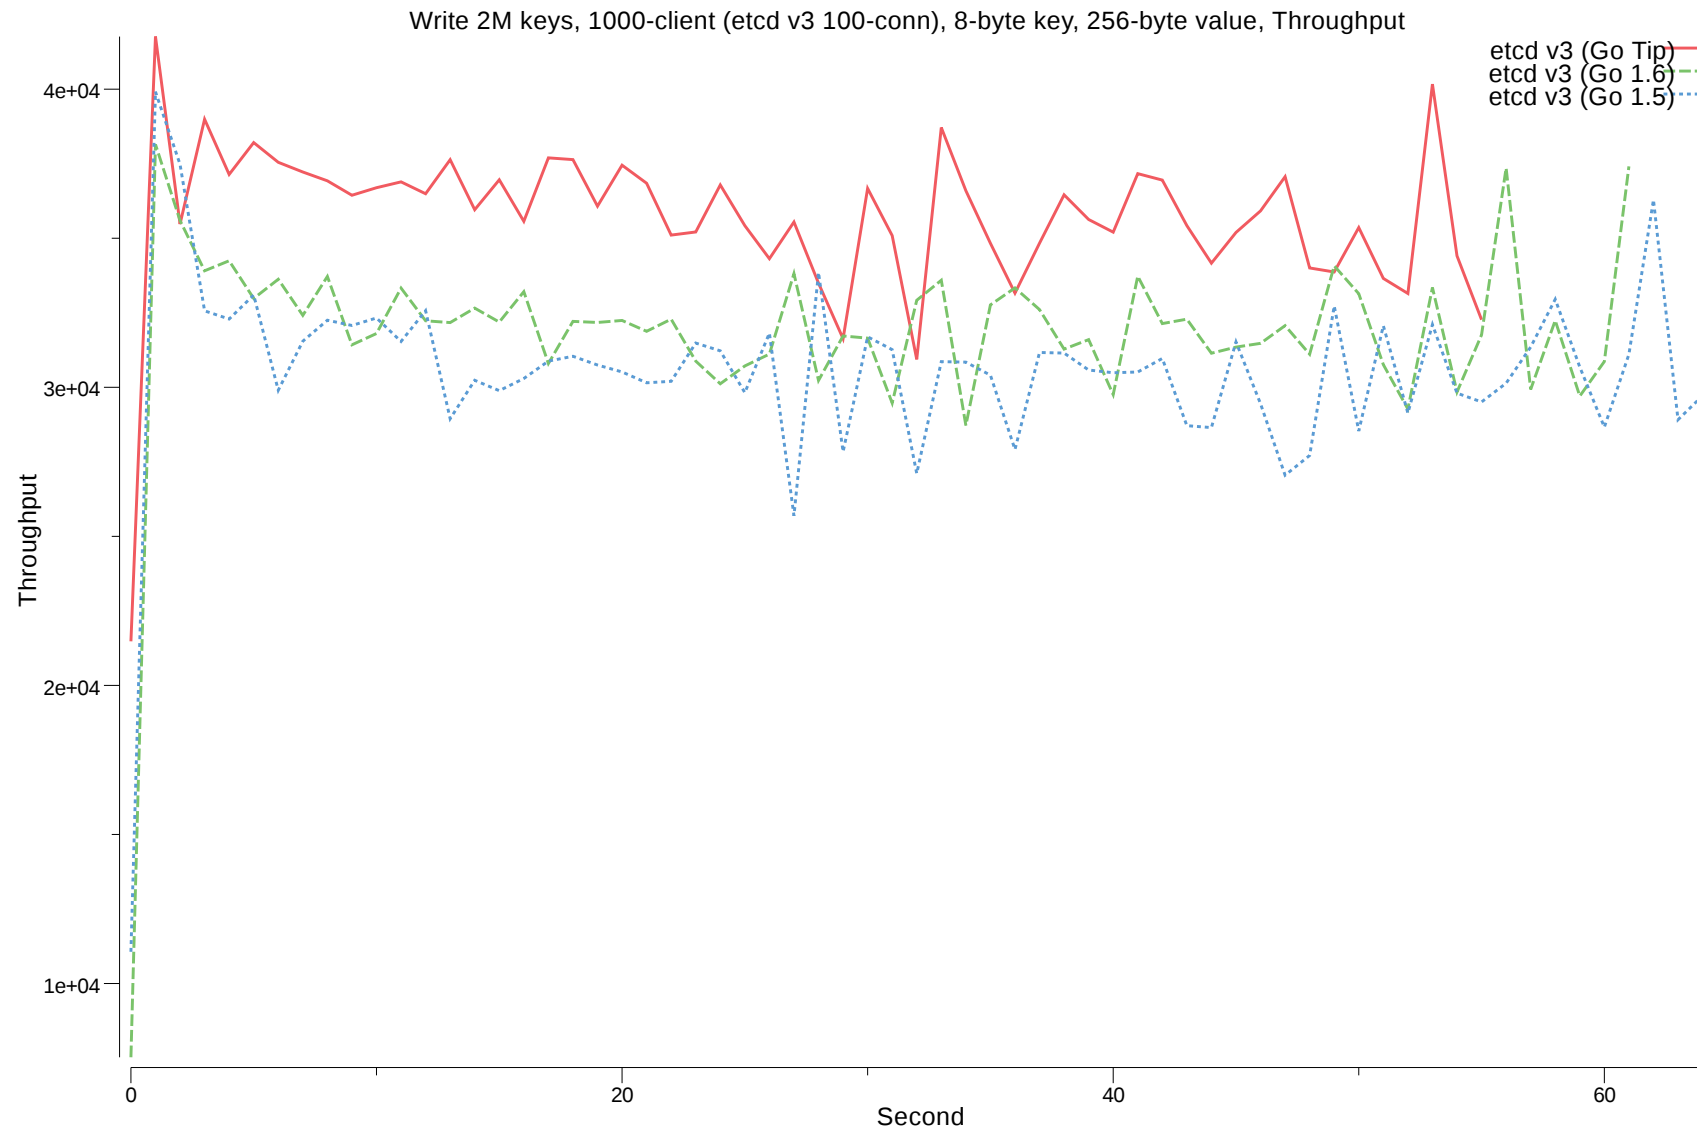
>12,000 failure injections per day

>2M injected for etcd v3

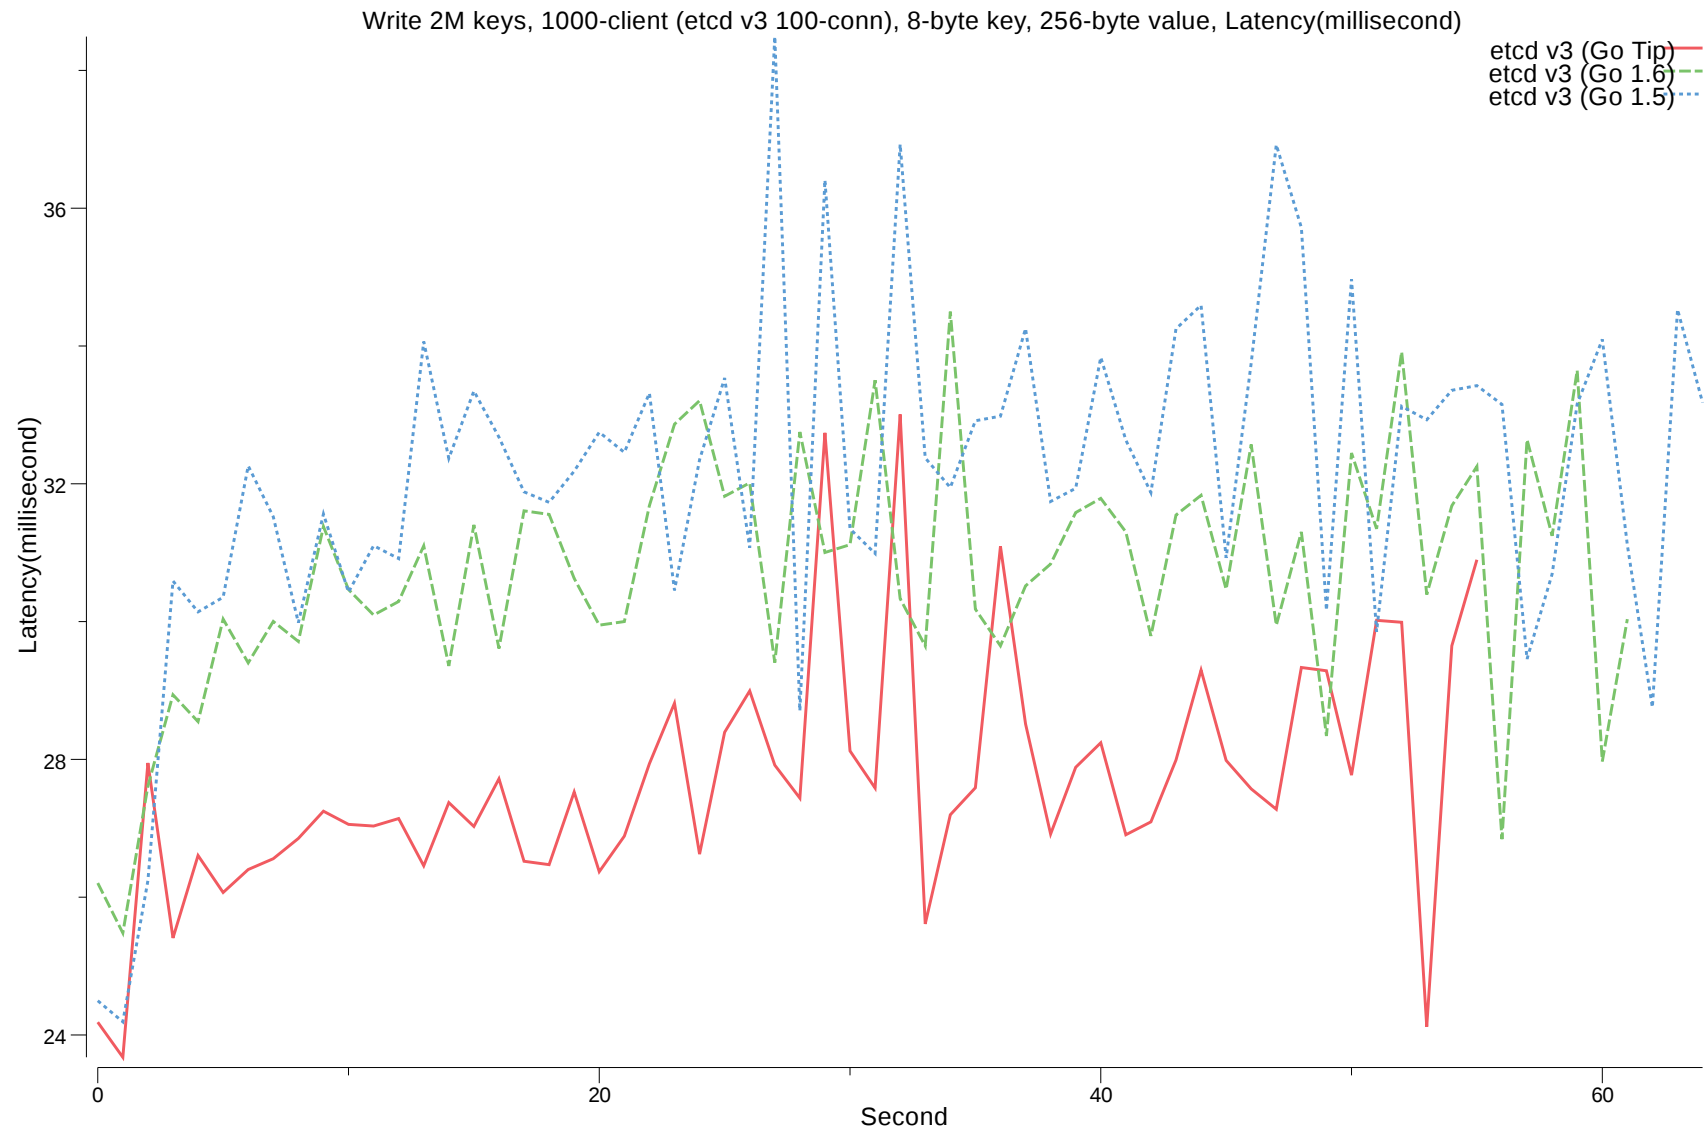
# Go 10 Tips

# #1 Use latest Go

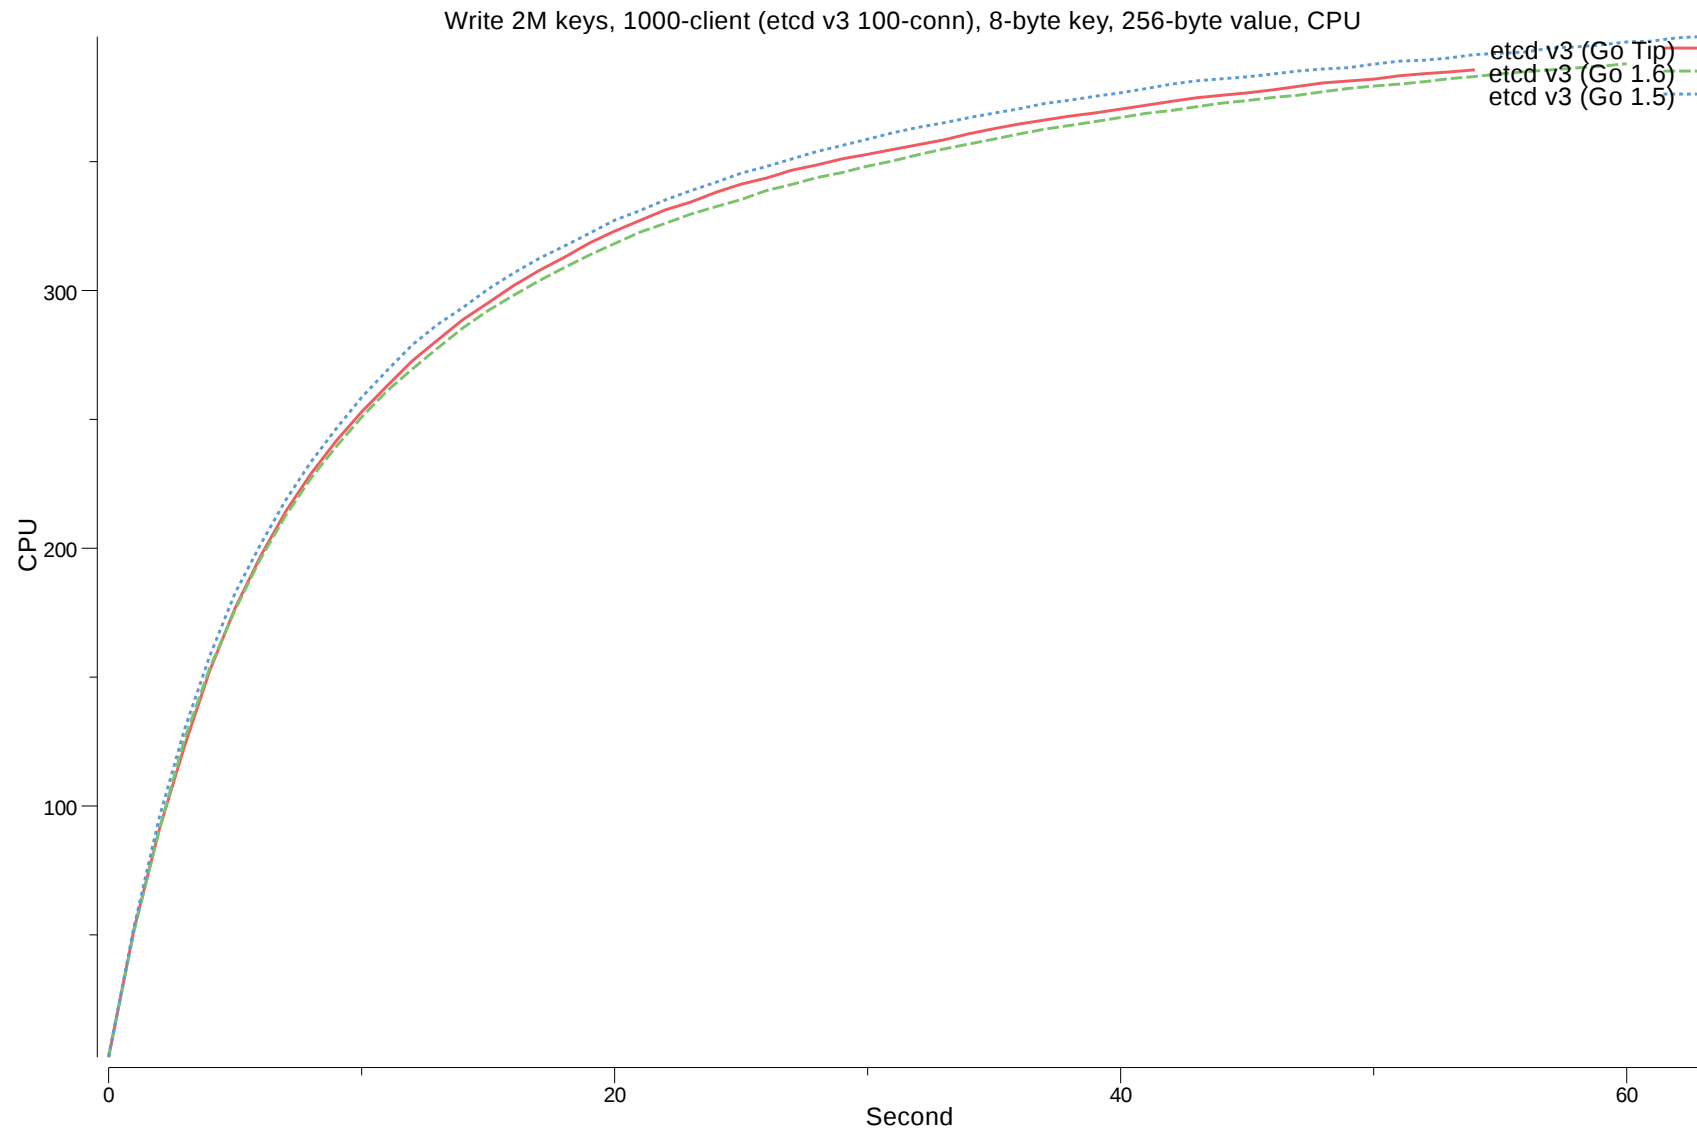
# Throughput




# Latency



# CPU



# Memory





## #2 Check slice allocation

Go slice capacity decides how much memory to use in the backing array

More capacity means more allocation.


Don't allocate more than you need.

[GH5238](https://github.com/coreos/etcd/pull/5238) (<https://github.com/coreos/etcd/pull/5238>)

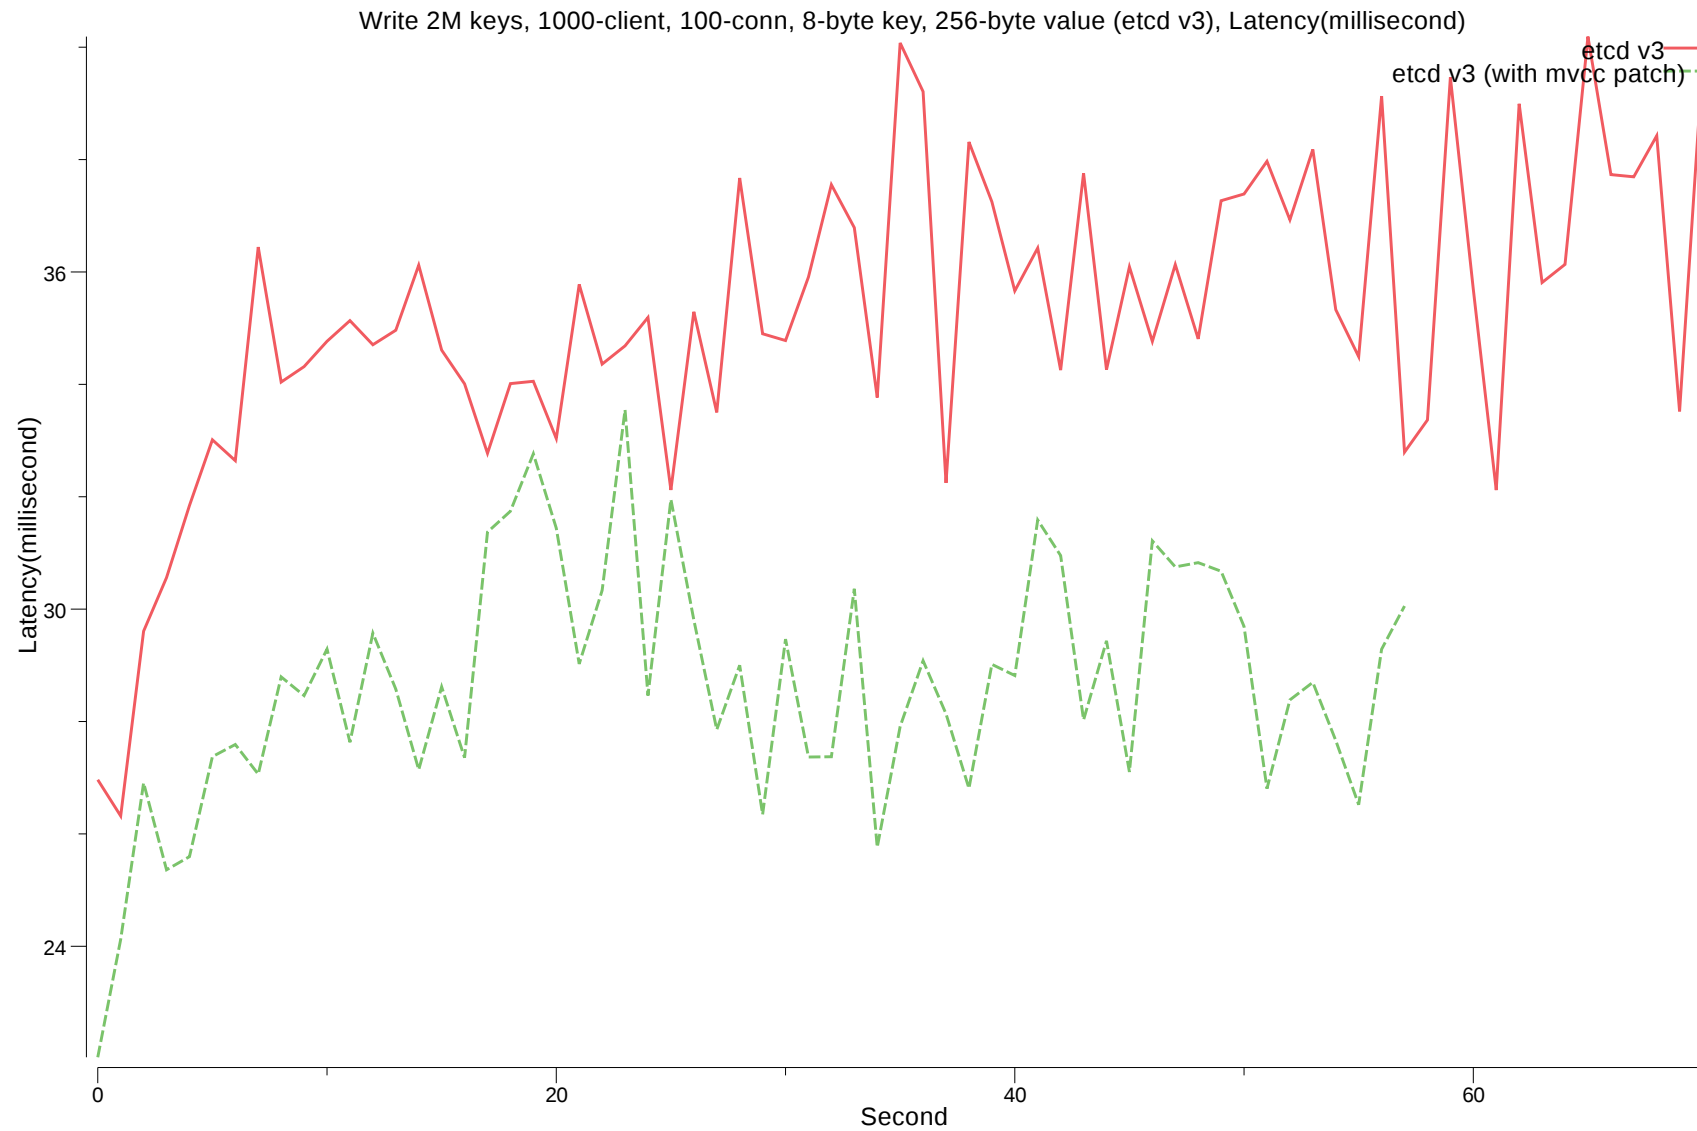
```
s.changes = make([]mvccpb.KeyValue, 0, 128)
s.changes = make([]mvccpb.KeyValue, 0, 4)    // better for etcd use case
```

Check your use case when making slice!

# Throughput



# Latency



## #3 Test your code

All with "go test" command

- unit tests
- integration tests
- functional tests
- benchmarks

Use expect package for e2e tests

```
proc, _ = expect.NewExpect("etcdctl", "get", "foo")  
_, err = proc.Expect("bar") // if err != nil, found a bug!
```

[github.com/coreos/etcd/pkg/expect](https://github.com/coreos/etcd/pkg/expect) (<https://godoc.org/github.com/coreos/etcd/pkg/expect>)

## #4 Check goroutine leaks

### Scan runtime.Stack

```
func TestMain(m *testing.M) {  
    v := m.Run()  
    if v == 0 && testutil.CheckLeakedGoroutine() {  
        os.Exit(1)  
    }  
    os.Exit(v)  
}  
  
func TestSample(t *testing.T) {  
    defer testutil.AfterTest(t)  
    ...  
}
```

- [net/http/main\\_test.go](https://github.com/golang/go/blob/master/src/net/http/main_test.go) ([https://github.com/golang/go/blob/master/src/net/http/main\\_test.go](https://github.com/golang/go/blob/master/src/net/http/main_test.go))
- [github.com/coreos/etcd/pkg/testutil](https://godoc.org/github.com/coreos/etcd/pkg/testutil) (<https://godoc.org/github.com/coreos/etcd/pkg/testutil>)

Highly recommend for projects with context.Context, gRPC

## #5 Always gofmt, go vet

### gofmt

```
Checking gofmt...
gofmt checking failed:
version/a.go
diff version/a.go gofmt/version/a.go
--- /tmp/gofmt6613415602016-05-15 04:07:11.087869561 +0000
+++ /tmp/gofmt2762292392016-05-15 04:07:11.087869561 +0000
@@ -15,5 +15,6 @@
 package version

 func myFunc() {
- a := 1
- a += 1 }
+     a := 1
+     a += 1
+}
```

### go vet

```
log.Fatalf("hello %d", "a")
// arg "a" for printf verb %d of wrong type: string
```

## #6 Write simple Go

- [github.com/dominikh/go-simple](https://github.com/dominikh/go-simple) (<https://github.com/dominikh/go-simple>) by Dominik
- Simplify your Go code

```
ok := true
if ok == true {} // X
if ok {}         // 0
```

Don't:

```
err := l.newStream()
if err != nil {
    return err
}
return nil
```

Do:

```
return l.newStream()
```

## #7 Check unused

- [github.com/dominikh/go-unused](https://github.com/dominikh/go-unused) (<https://github.com/dominikh/go-unused>) by Dominik
- Finds unused constants, variables, functions and types

```
func reportMetrics() {}  
// func reportMetrics is unused
```

Found bugs in etcd [GH4955](https://github.com/coreos/etcd/pull/4995/files) (<https://github.com/coreos/etcd/pull/4995/files>)



## #8 Use goword

- [github.com/chzchzchz/goword](https://github.com/chzchzchz/goword) (https://github.com/chzchzchz/goword) by Anthony (etcd team)

### Comment checker

```
// This.  
func Hello() {} // This. (godoc-export: This -> Hello?)"
```

### Spell checker

```
// Hello retuens.  
func Hello() {} // Hello retuens. (spell: retuens -> returns?)
```

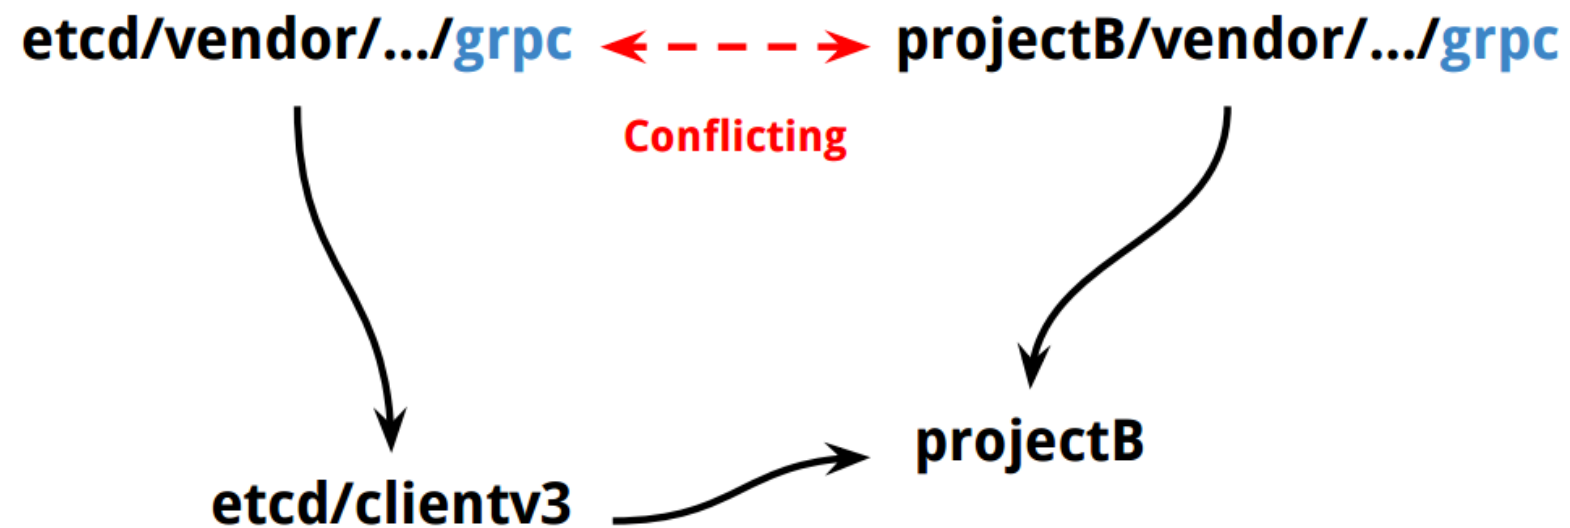
goword understands your Go syntax

## #9 Document with godoc

- etcd must be easy to use
- etcd needs good documentation
- Use godoc to document our client package with code examples
- [etcd/clientv3](https://github.com/coreos/etcd/tree/master/clientv3) (<https://github.com/coreos/etcd/tree/master/clientv3>)

## #10 vendor

etcd use case is complicated...



## #10 vendor

### Problem

- etcd "client" package is used within etcd repo (etcdctl)
- etcd "client" imports "grpc" and vendors it
- Project B imports this etcd "client" package
- Project B also uses "grpc", but from different import path

Now two projects has conflicting "grpc" code [GH566](https://github.com/grpc/grpc-go/issues/566) (<https://github.com/grpc/grpc-go/issues/566>)

```
panic: http: multiple registrations for /debug/requests
```

## #10 vendor

Solution [GH4950](https://github.com/coreos/etcd/pull/4950) (<https://github.com/coreos/etcd/pull/4950>)

- Create symlinks inside cmd directory
- In -s main.go cmd/main.go
- cmd/vendor

To update dependency

```
ln -s cmd/vendor vendor && godep save
```


## #10 vendor

Still go-get-able.

No conflicts with other projects.

Works, even on Windows!

```
clientv3
etcdserver
mvcc
vendor
cmd
  vendor
  ../main.go@
main.go
```



# Thank you

Gyu-Ho Lee

CoreOS

[gyu\\_ho.lee@coreos.com](mailto:gyu_ho.lee@coreos.com) (mailto:gyu\_ho.lee@coreos.com)

<https://github.com/coreos/etcd> (https://github.com/coreos/etcd)

