

etcd: mission critical key-value store

Gopherfest
16 May 2016

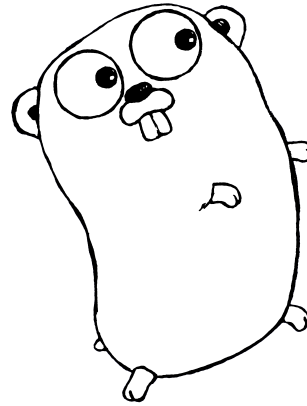
Gyu-Ho Lee
CoreOS

Welcome

Slides are here:

github.com/gyuho/presentations (<https://github.com/gyuho/presentations>)

Go + etcd



Agenda

- What is etcd
- Why
- Go
- Q/A

What is etcd

etcd is ...

- Distributed key-value store
- Open source github.com/coreos/etcd (<https://github.com/coreos/etcd>)
- Since June 2013
- Still new (vs ZooKeeper since May 2008)
- Google, Red Hat, EMC, Cisco, Huawei, Baidu, Alibaba...

etcd API

```
cli.Put(ctx, "foo", "bar", Lease)
cli.Get(ctx, "foo")
cli.Delete(ctx, "foo")

// Transaction
kvc.Txn(ctx).
  If(clientv3.Compare(clientv3.Value("key"), ">", "abc")). // txn value comparisons are lexical
  Then(clientv3.OpPut("key", "XYZ")).                      // this runs, since 'xyz' > 'abc'
  Else(clientv3.OpPut("key", "ABC")).
  Commit()

// Watch for updates on key
ch := cli.Watch(ctx, "foo")
for res := range ch {}

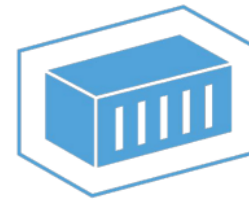
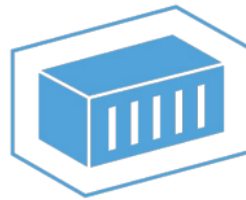
// Distributed locks
mu := concurrency.NewMutex(cli, "foo")
mu.Lock()
mu.Unlock()
```

Why etcd

Use etcd to store configuration

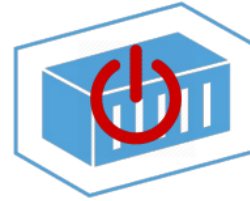
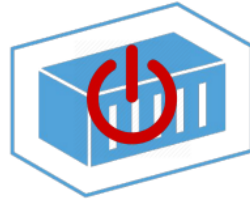
Updates

Security updates?



How would you update the cluster of machines?

Traditional way



- Reboot with downtime
- Manual

CoreOS updates with etcd

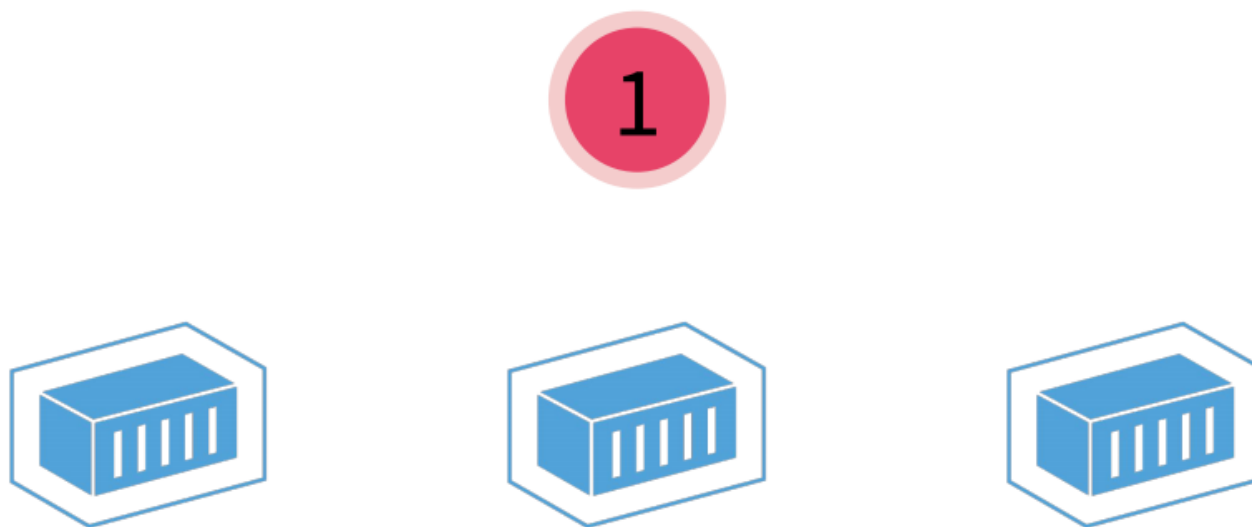
1032.1.0 Release Date: May 5, 2016 kernel: 4.5.2 rkt: 1.2.1 docker: 1.10.3 etcd: 0.4.9, 2.3.2 fleet: 0.11.7 systemd: 229

Security Updates:

- OpenSSL [1.0.2h](#) for [CVE-2016-2105](#), [CVE-2016-2106](#), [CVE-2016-2107](#), [CVE-2016-2109](#), [CVE-2016-2176](#)
- ntpd [4.2.8p7](#) for [CVE-2016-1551](#), [CVE-2016-1549](#), [CVE-2016-2516](#), [CVE-2016-2517](#), [CVE-2016-2518](#), [CVE-2016-2519](#), [CVE-2016-1547](#), [CVE-2016-1548](#), [CVE-2015-7704](#), [CVE-2015-8138](#), [CVE-2016-1550](#)
- git [2.7.3-r1](#) for [CVE-2015-7545](#), [CVE-2016-2315](#), [CVE-2016-2315](#)
- jq [1.5-r2](#) for [CVE-2015-8863](#)

CoreOS updates with etcd

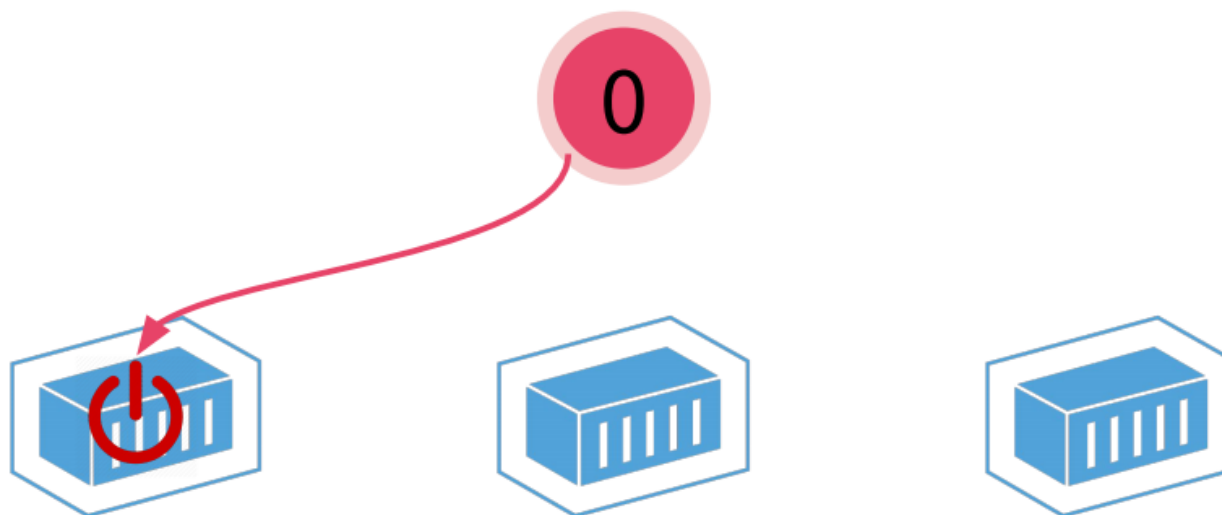
- CoreOS is an open-source Linux OS
- Automatic, No-downtime updates with etcd



Powered by etcd and github.com/coreos/locksmith (<https://github.com/coreos/locksmith>)

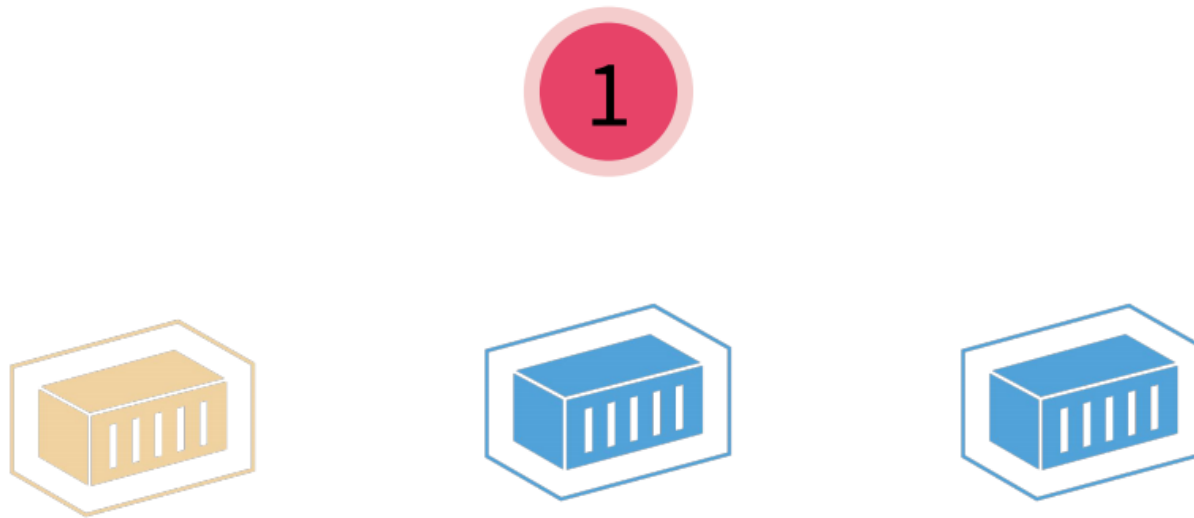
CoreOS updates with etcd

- locksmith stores semaphore values in etcd
- ensure that only a subset of a cluster of machines are rebooting at any given time

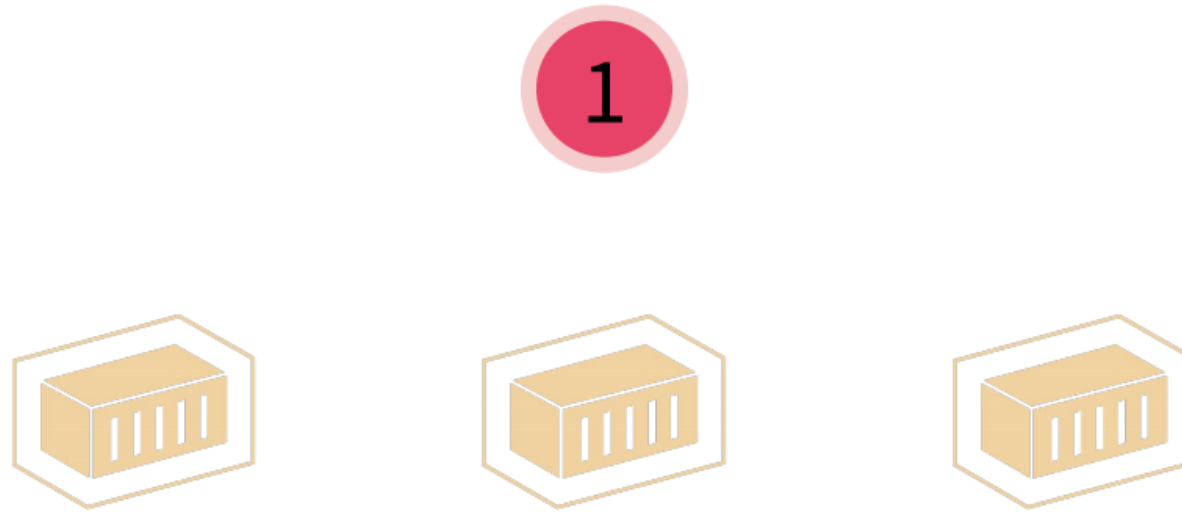


CoreOS updates with etcd

- Automatic
- No downtime



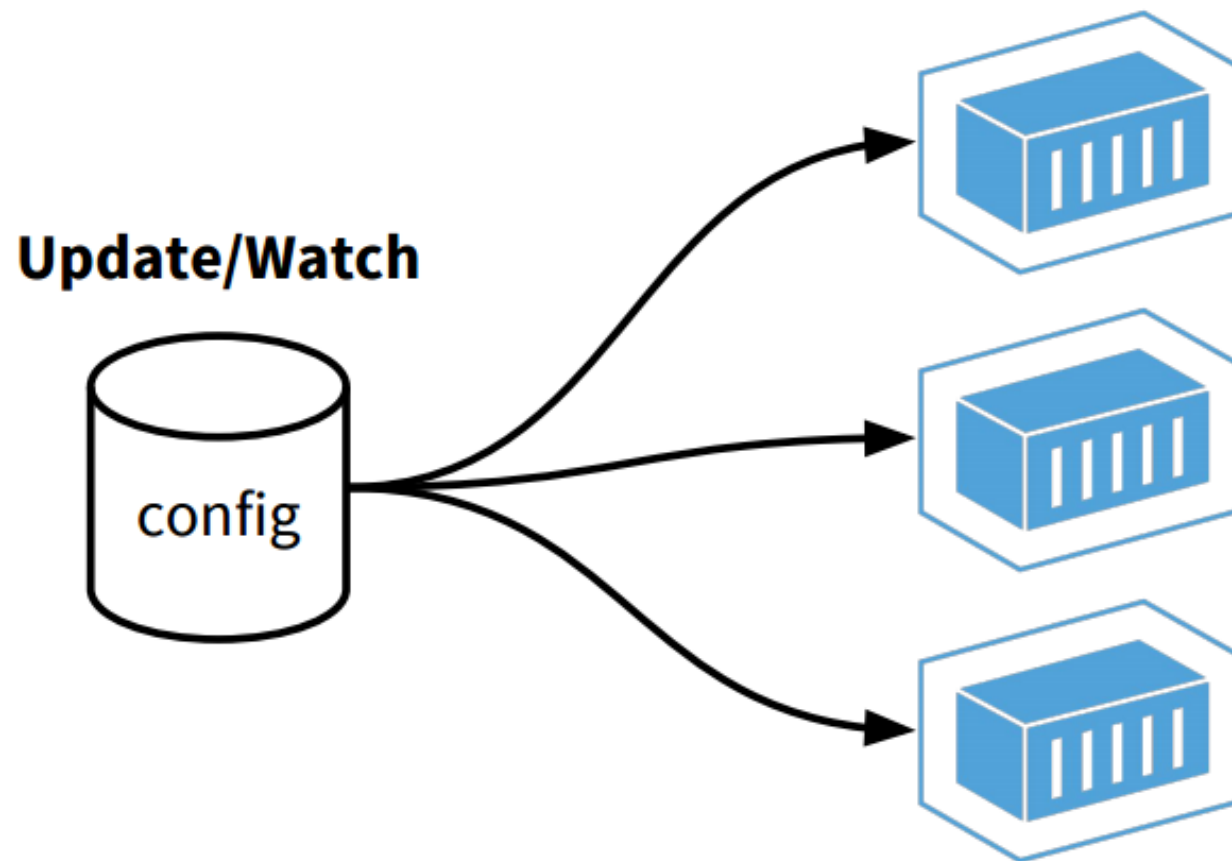
CoreOS updates with etcd



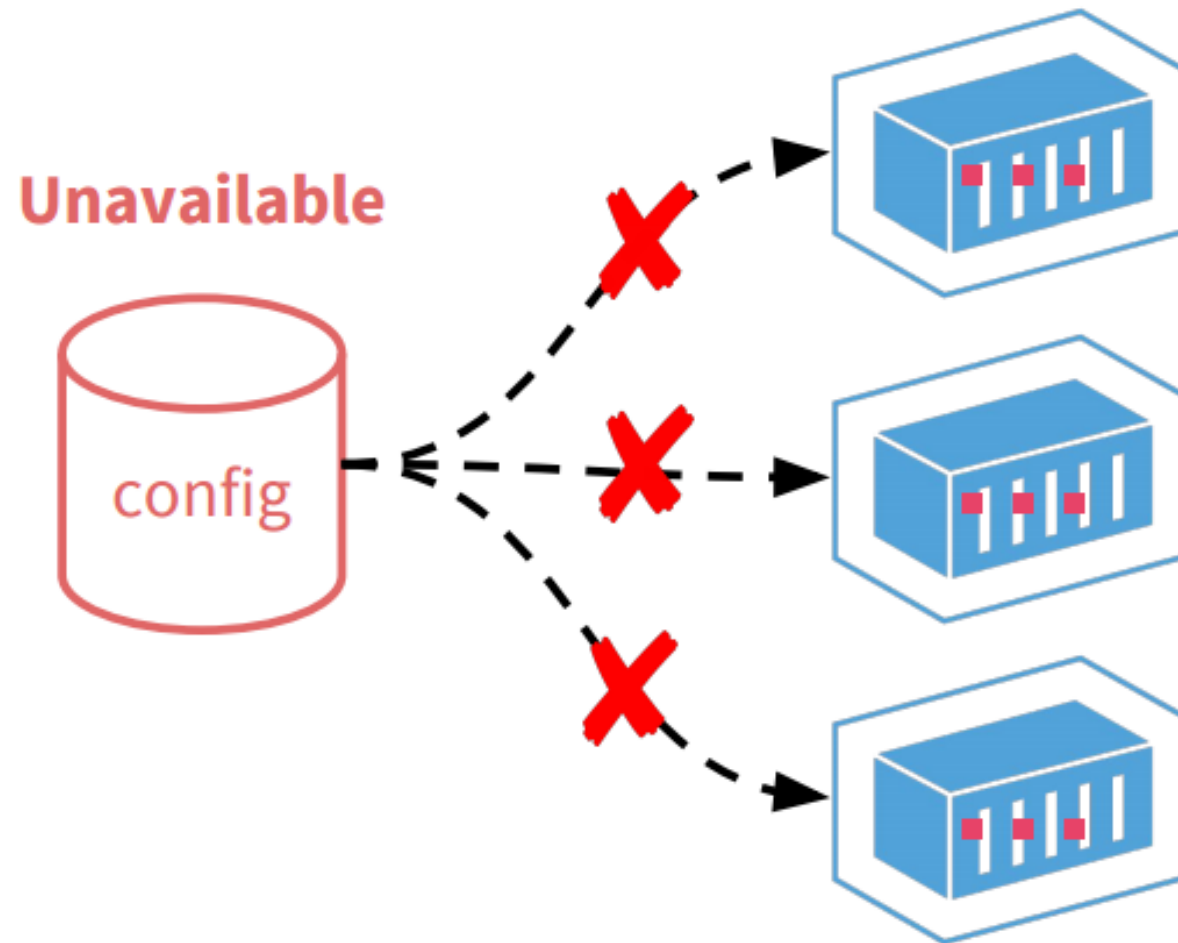
Your cluster is now secured

Use etcd for "critical" configuration

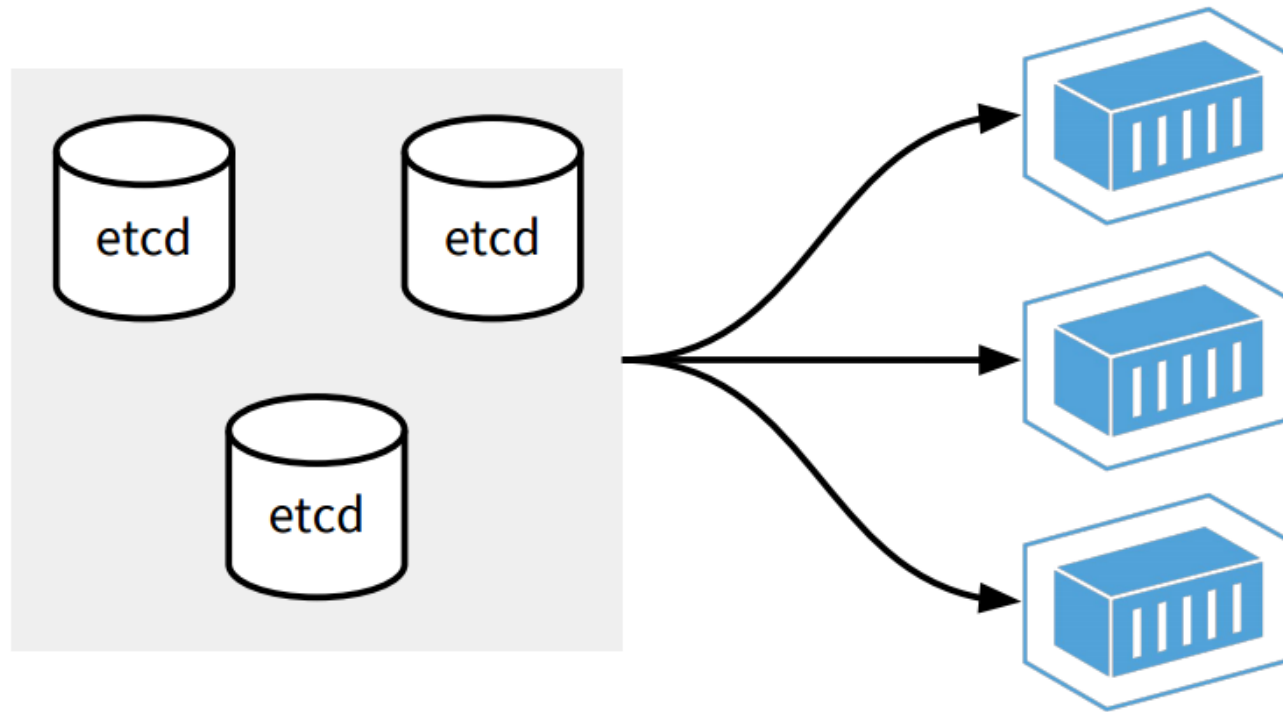
Bad practice



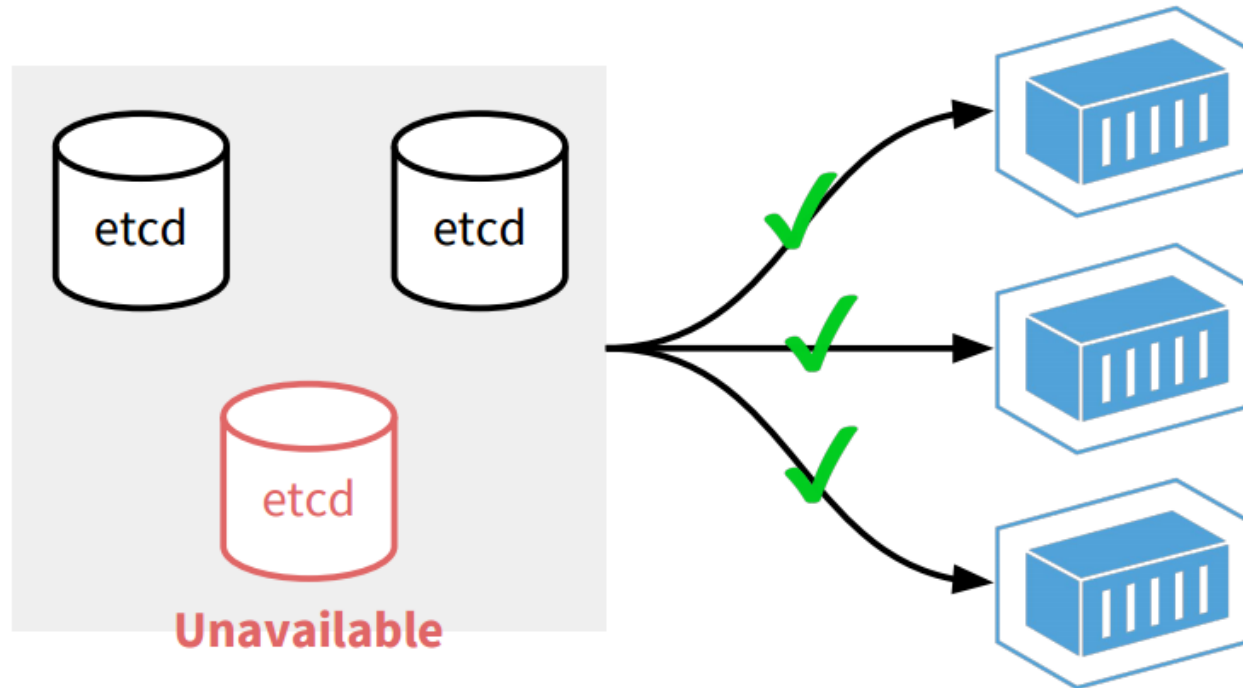
Bad practice



Good practice



Good practice



Demo

play.etcd.io (<http://play.etcd.io>)

Join me!

etcd

Consistent view of critical configuration

- Strong consistency (no stale reads)
- Different than eventual consistency (conflicts, latest timestamps wins)

Highly available configuration store

- Resilient to a single point of failures & network partitions

Watchable

- Push configuration updates to application

Why not ZooKeeper or Consul?

etcd Project Status

etcd Performance

etcd v3

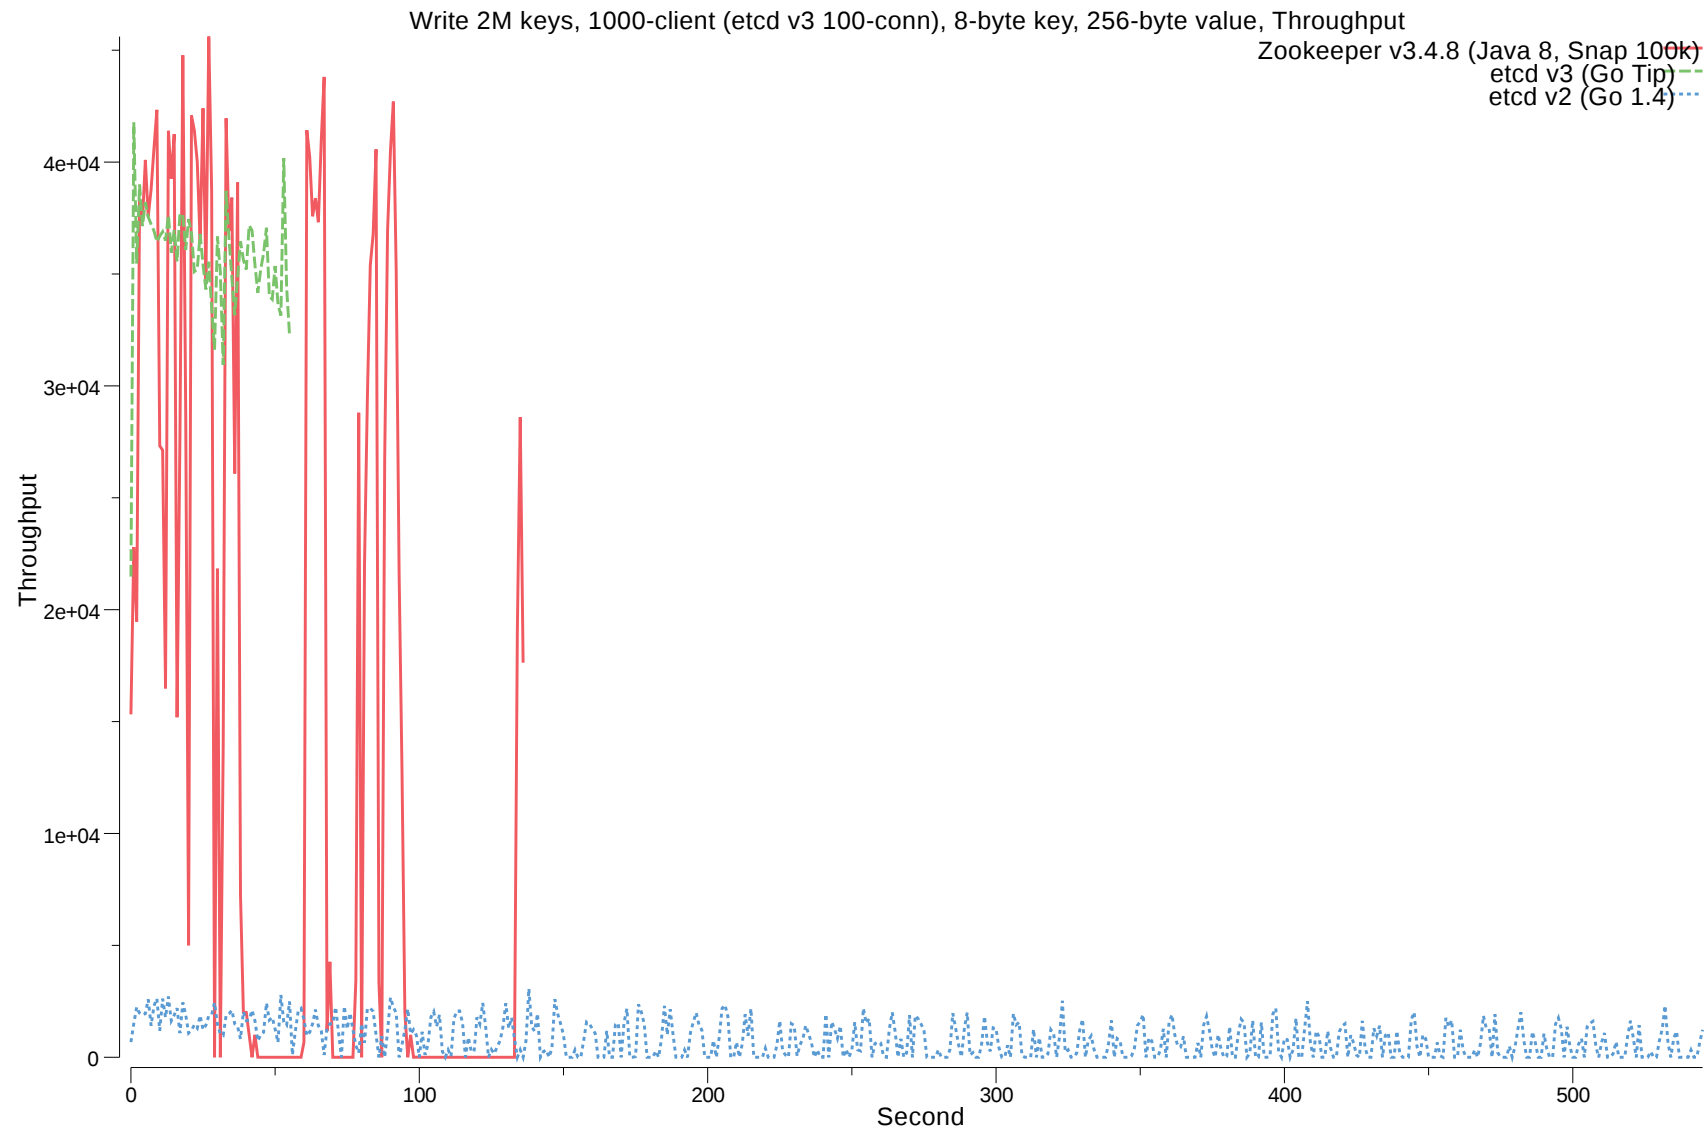
BoltDB (<https://github.com/boltdb/bolt>)

- B+tree disk storage
- Incremental snapshot
- vs ZooKeeper snapCount 10,000

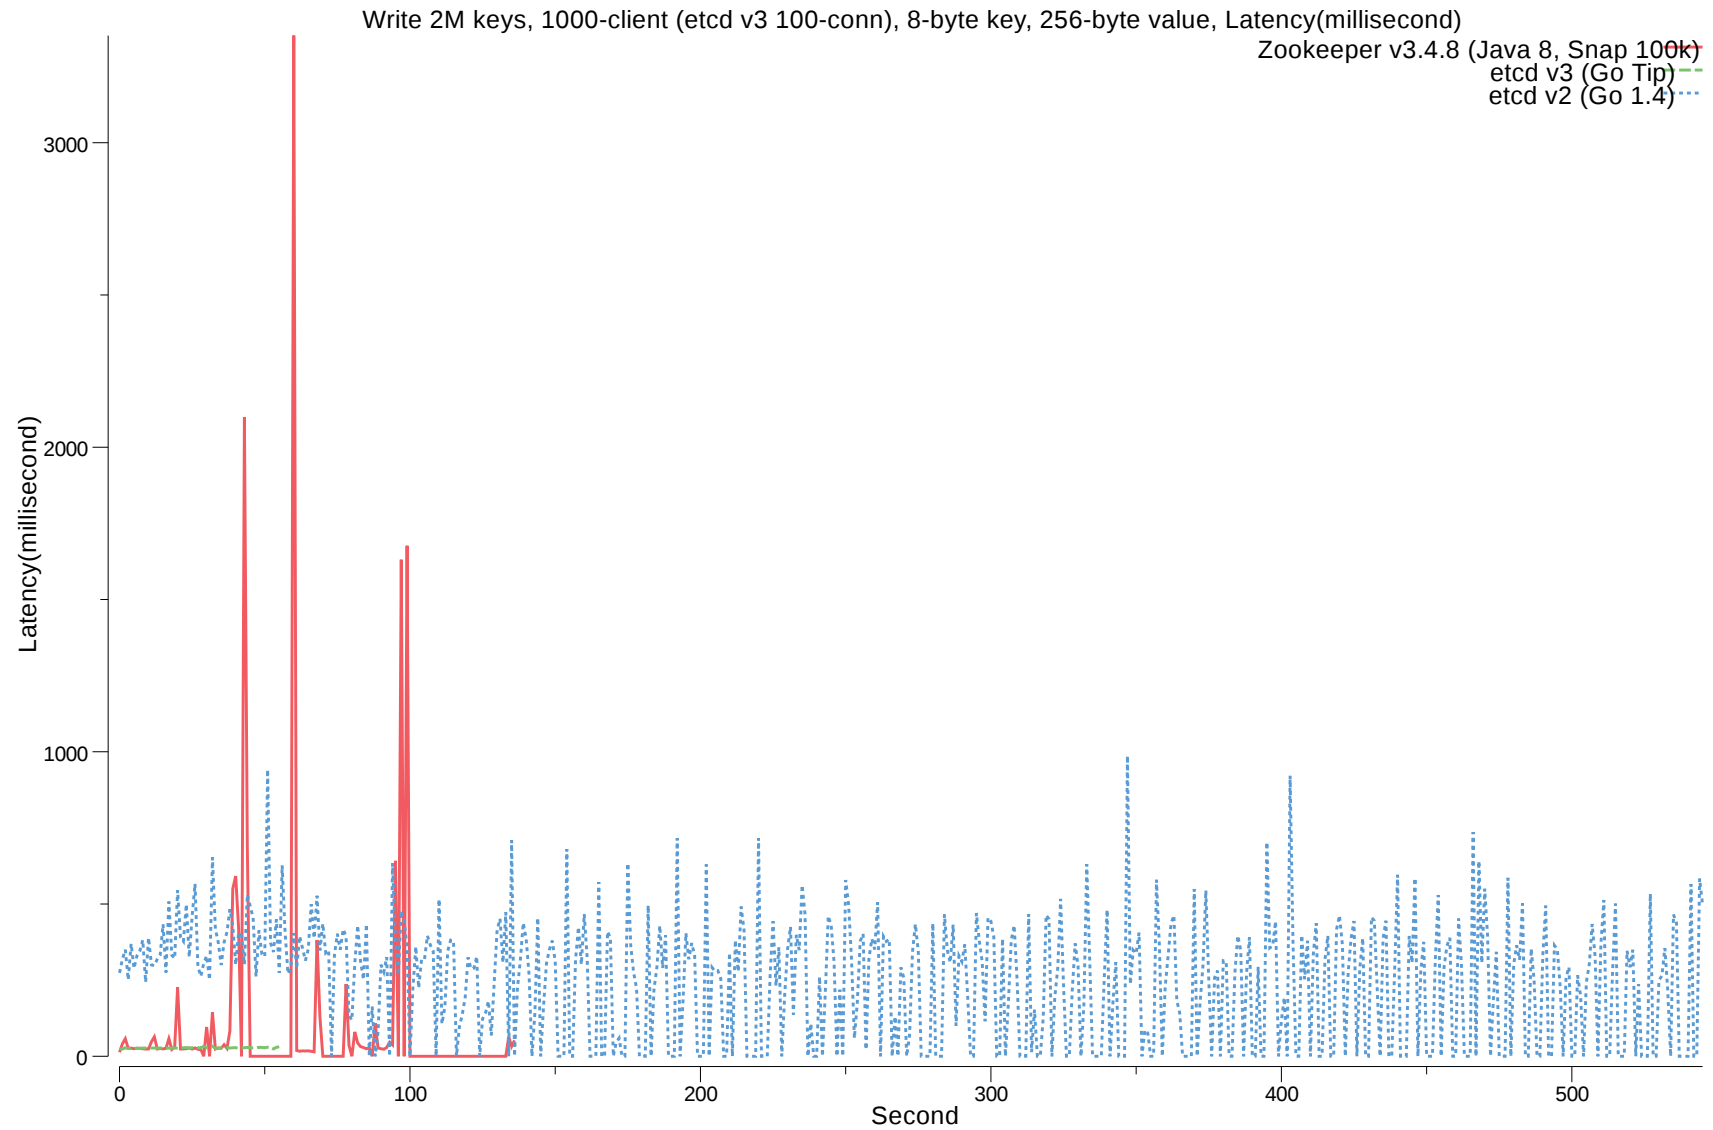
gRPC (<http://www.grpc.io/>)

- Protocol Buffer
- HTTP/2
- streams, less TCP congestions

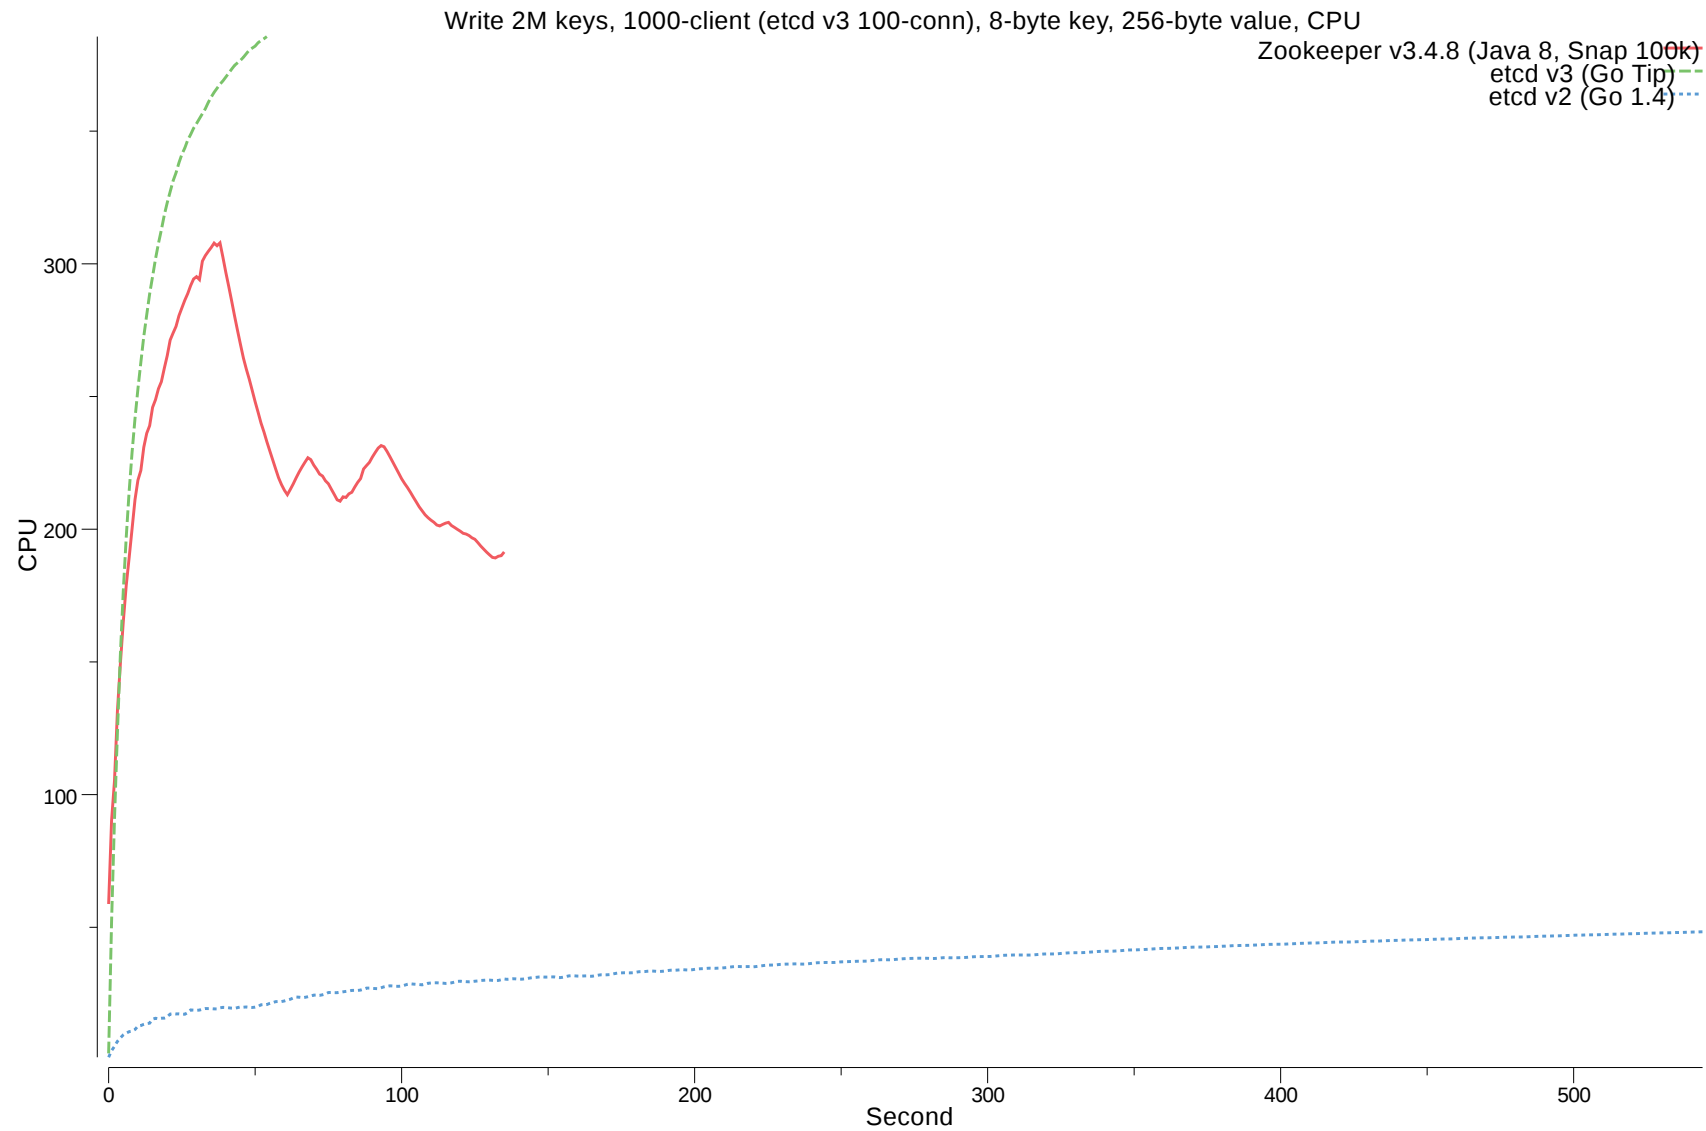
Throughput



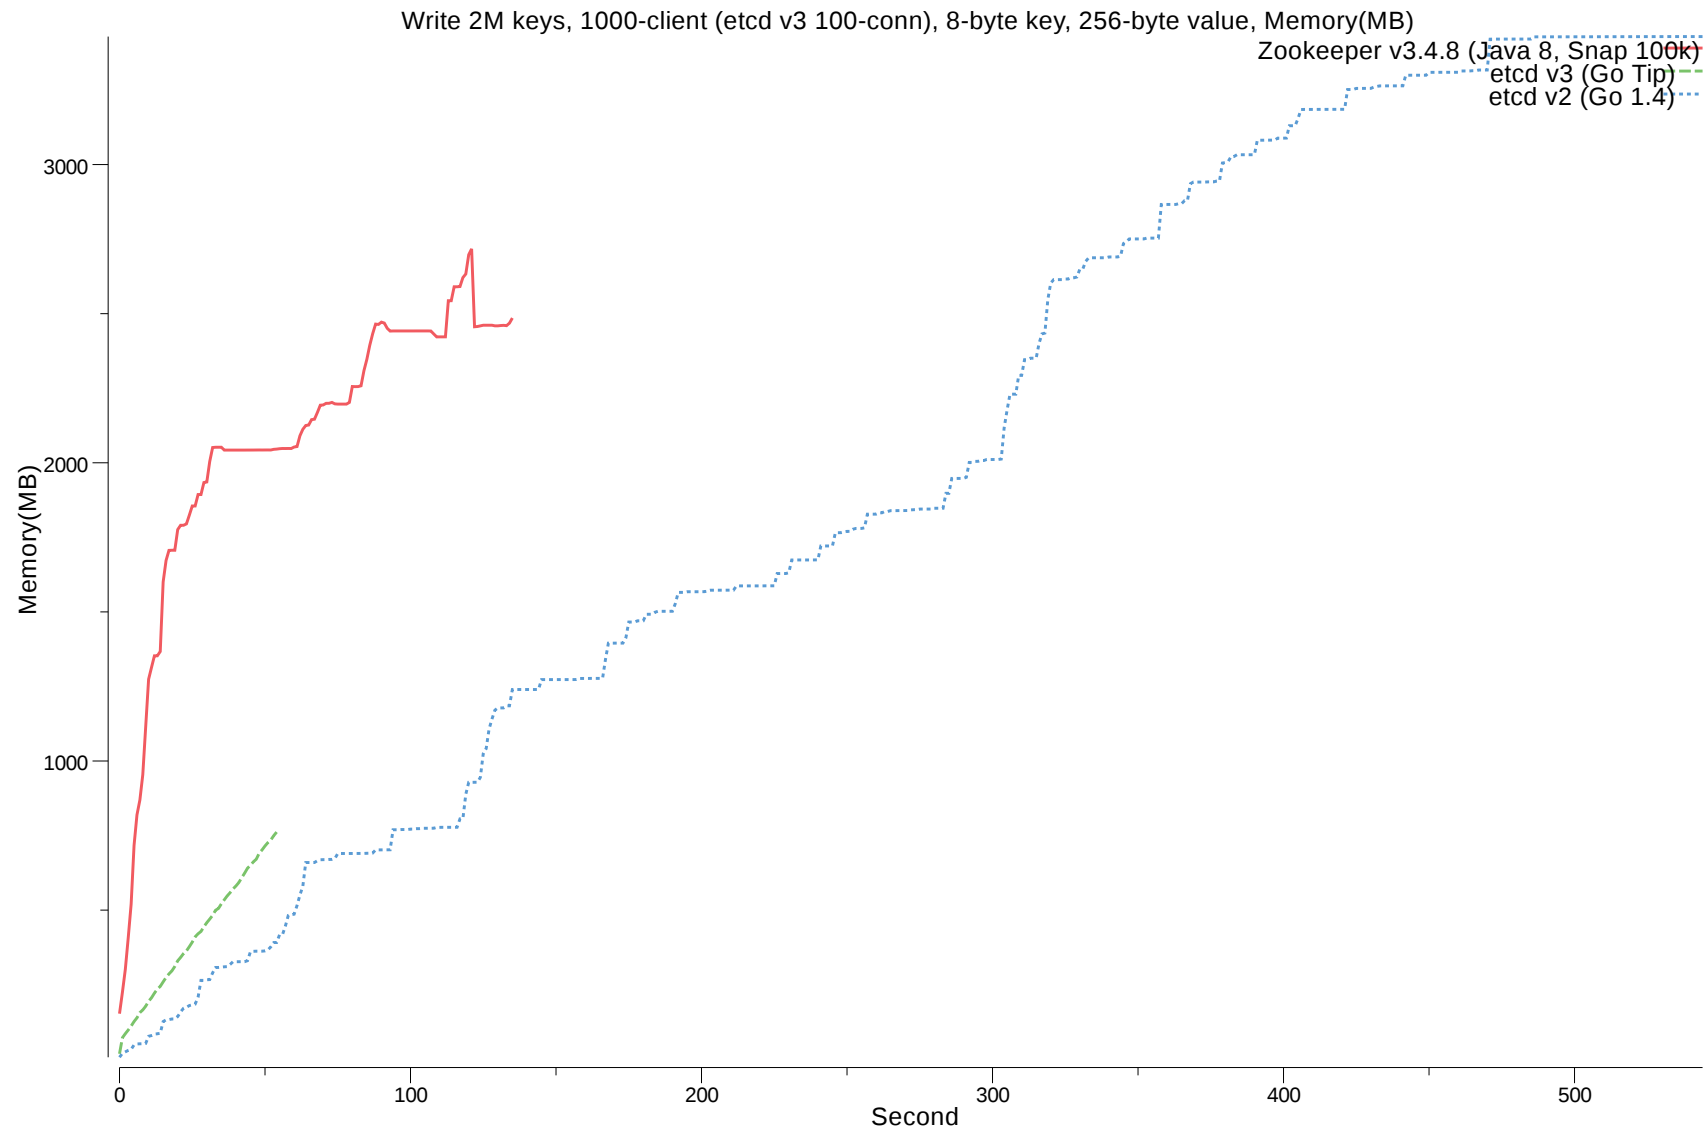
Latency



CPU

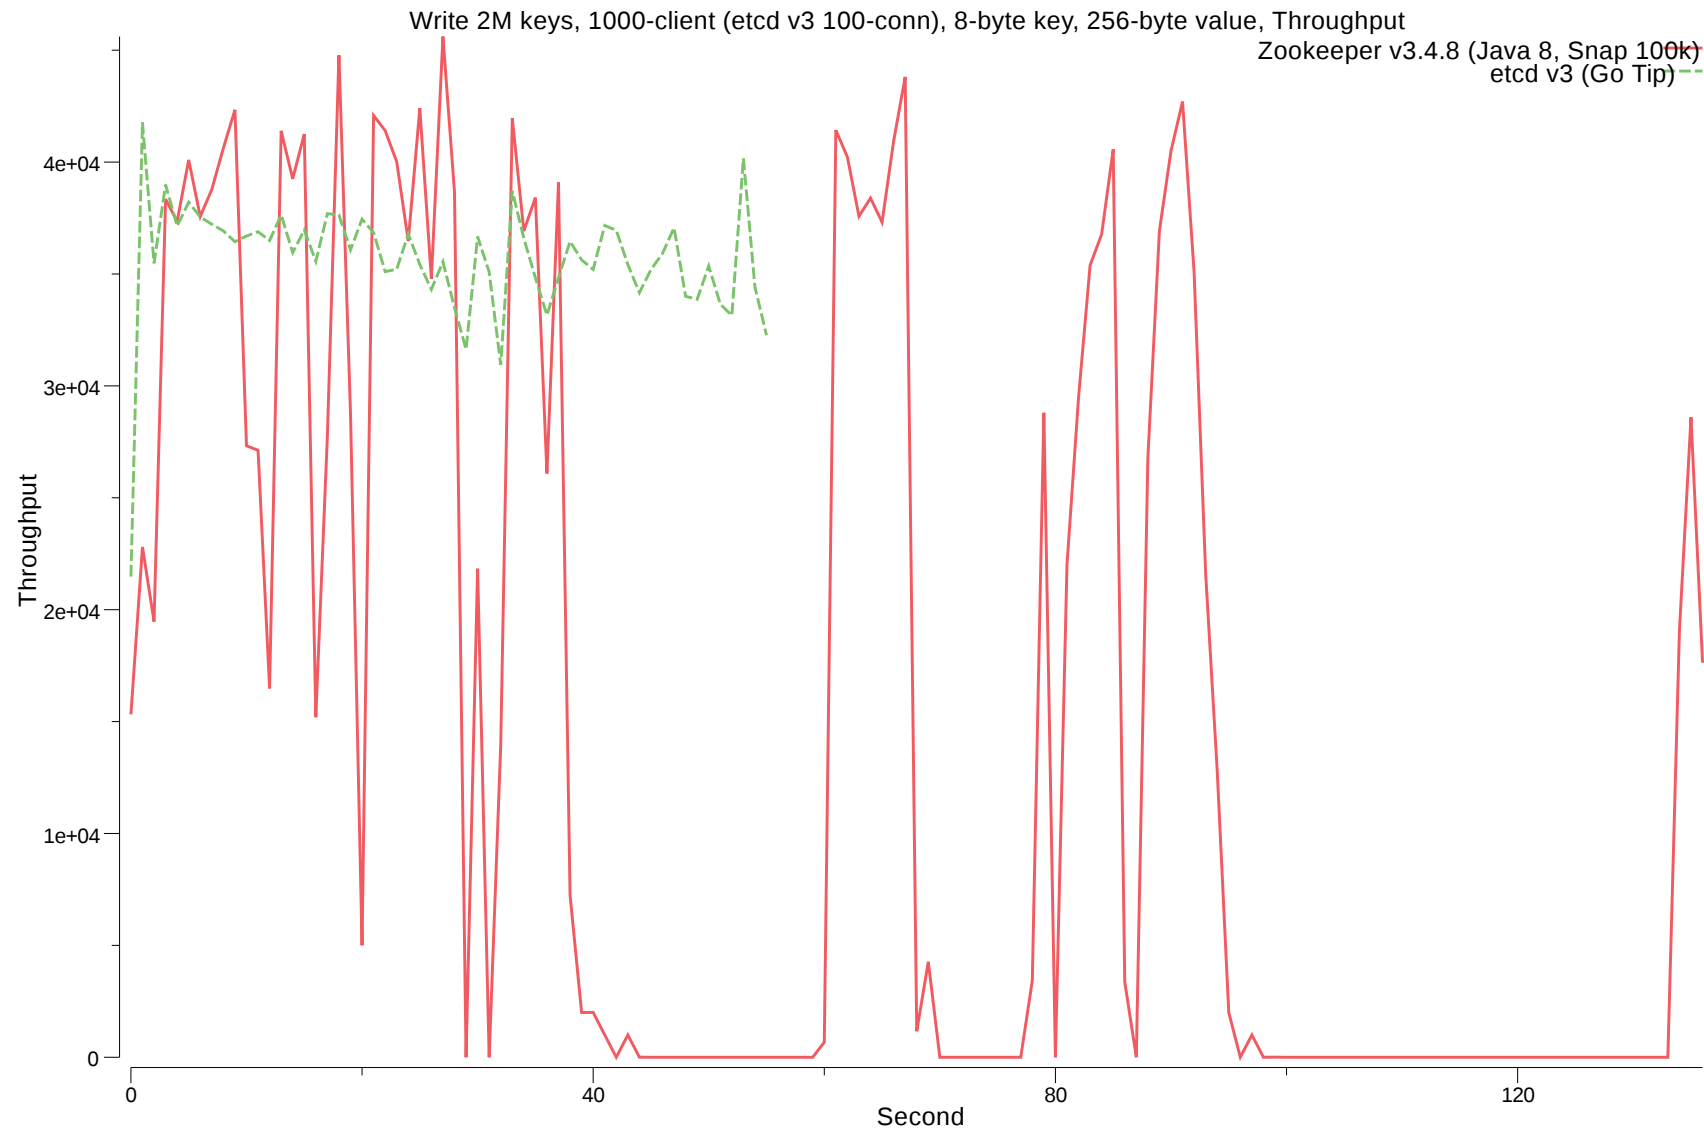


Memory

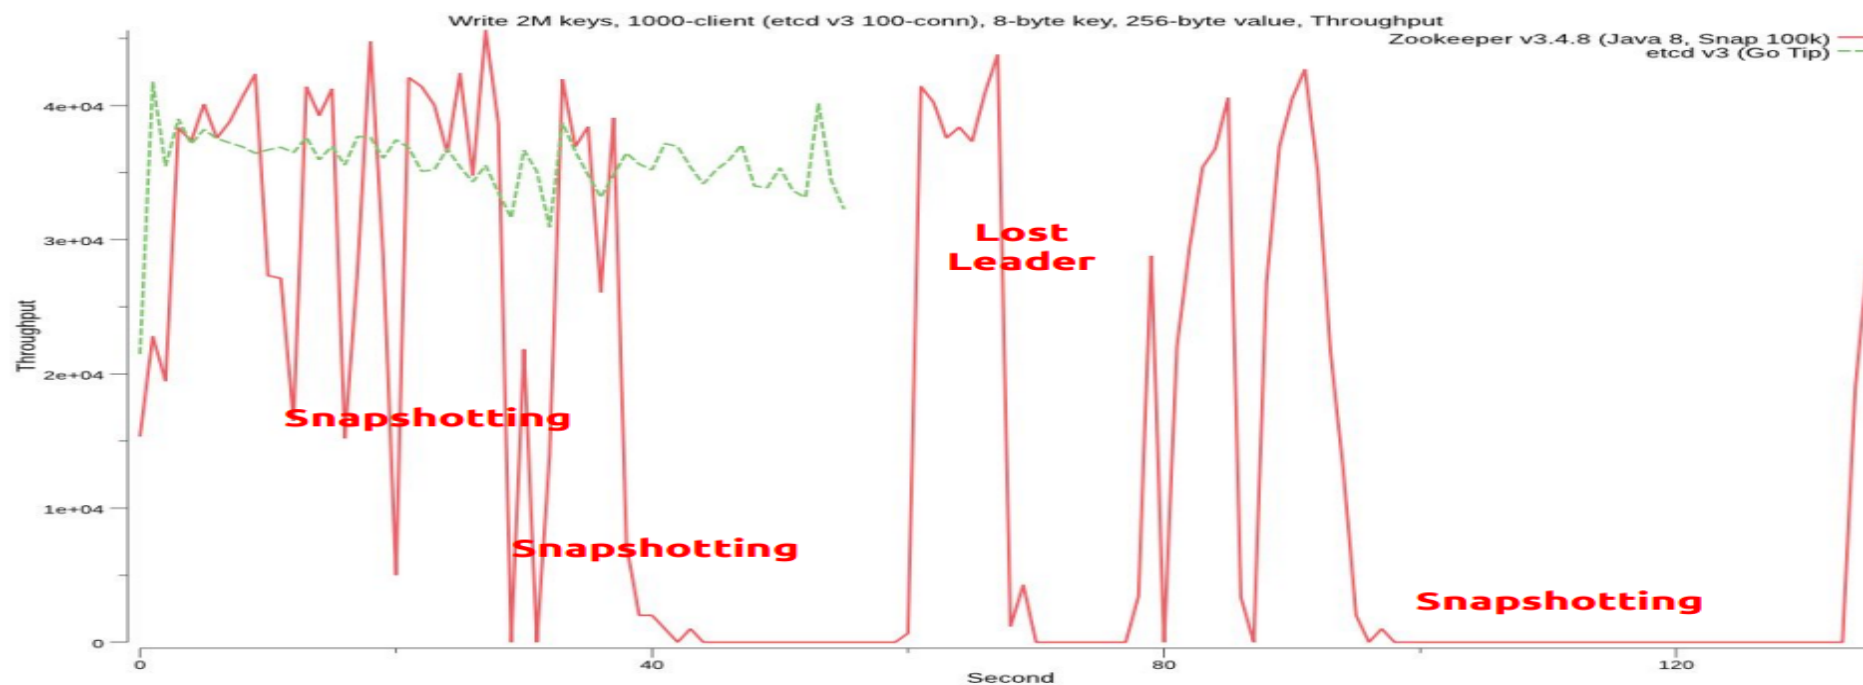


etcd Reliability

Throughput



etcd Reliability



ZooKeeper logs

```
07:16:35 [Snapshot Thread:FileTxnSnapLog@240] - Snapshotting...
07:16:43 fsync-ing the write ahead log in SyncThread:3 took 1224ms...
07:16:46 fsync-ing the write ahead log in SyncThread:3 took 3205ms... // Snapshotting
...
07:17:14 [FastLeaderElection@818] - New election... // Leader Election
```

etcd Reliability

We test...

- Kill all members
- Kill majority of members
- Kill leader
- Network partition
- Network latency
- More...

etcd Reliability

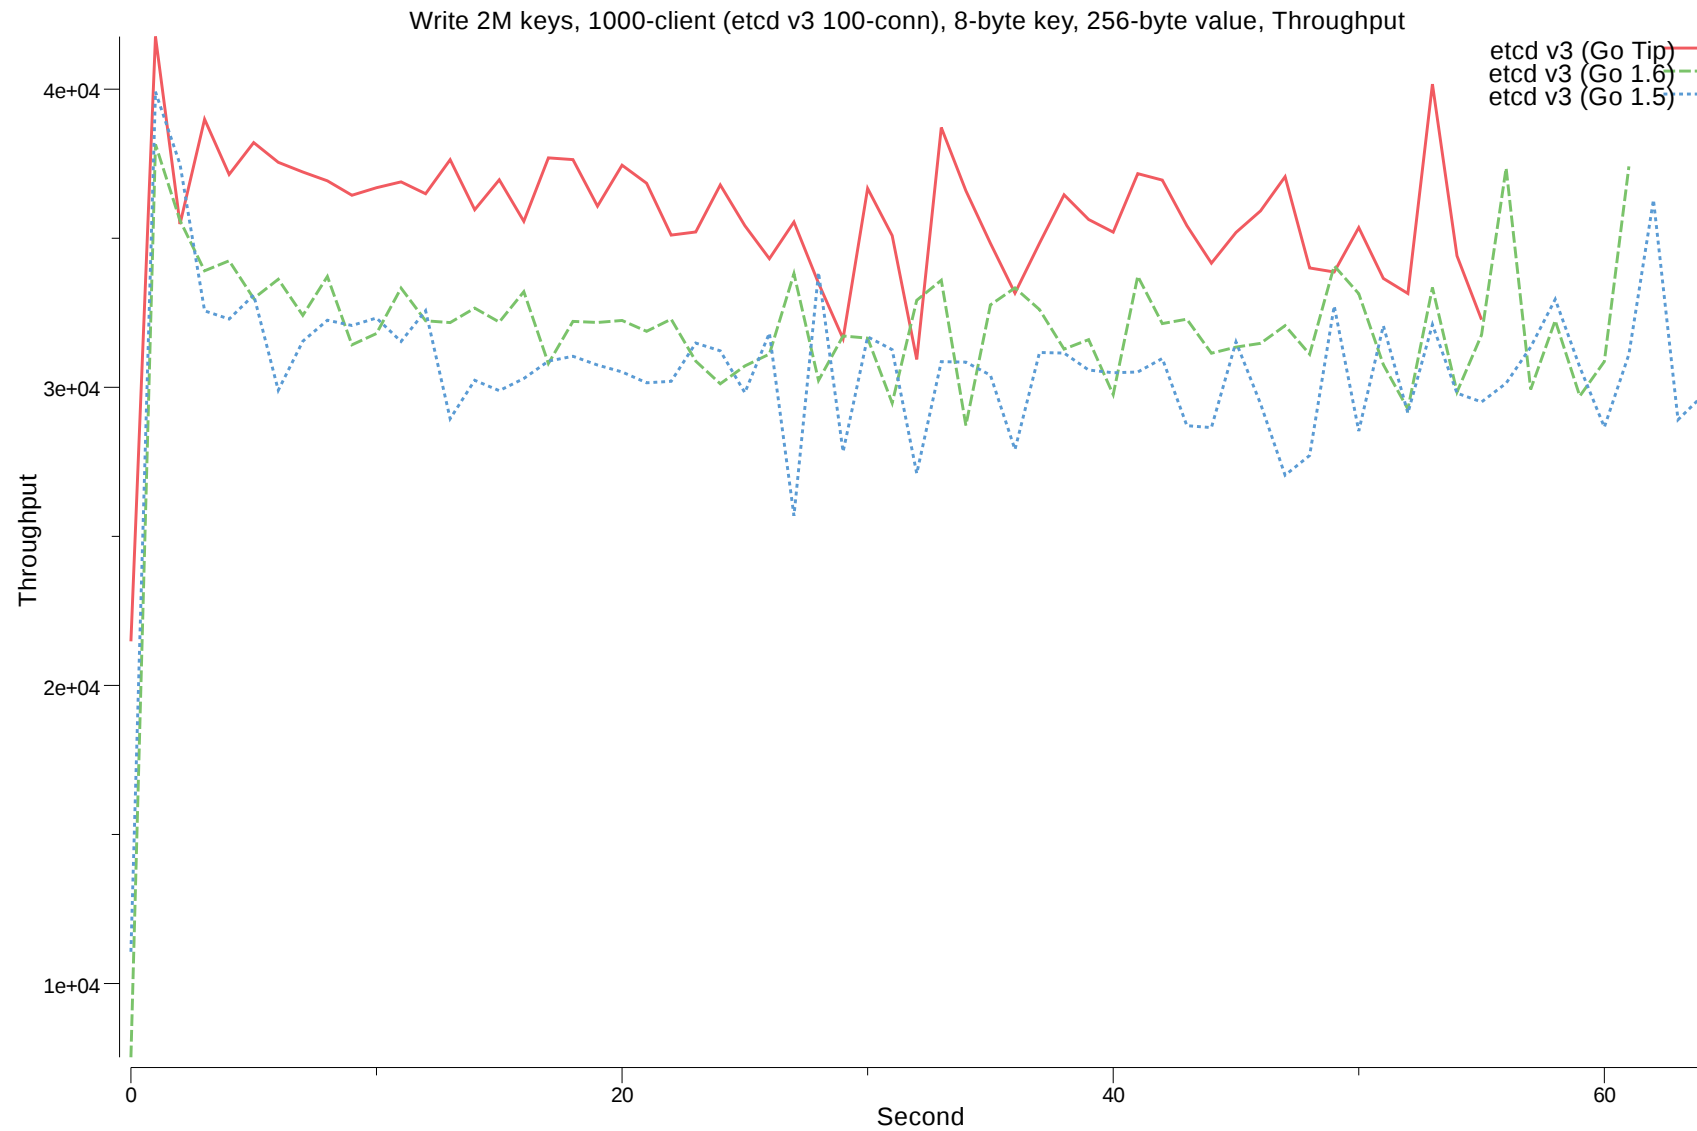
Intensive failure injection testing

- dash.etcd.io (<http://dash.etcd.io/dashboard/db/functional-tests>)
- >12,000 failure injections per day
- >1.5M injected for etcd v3

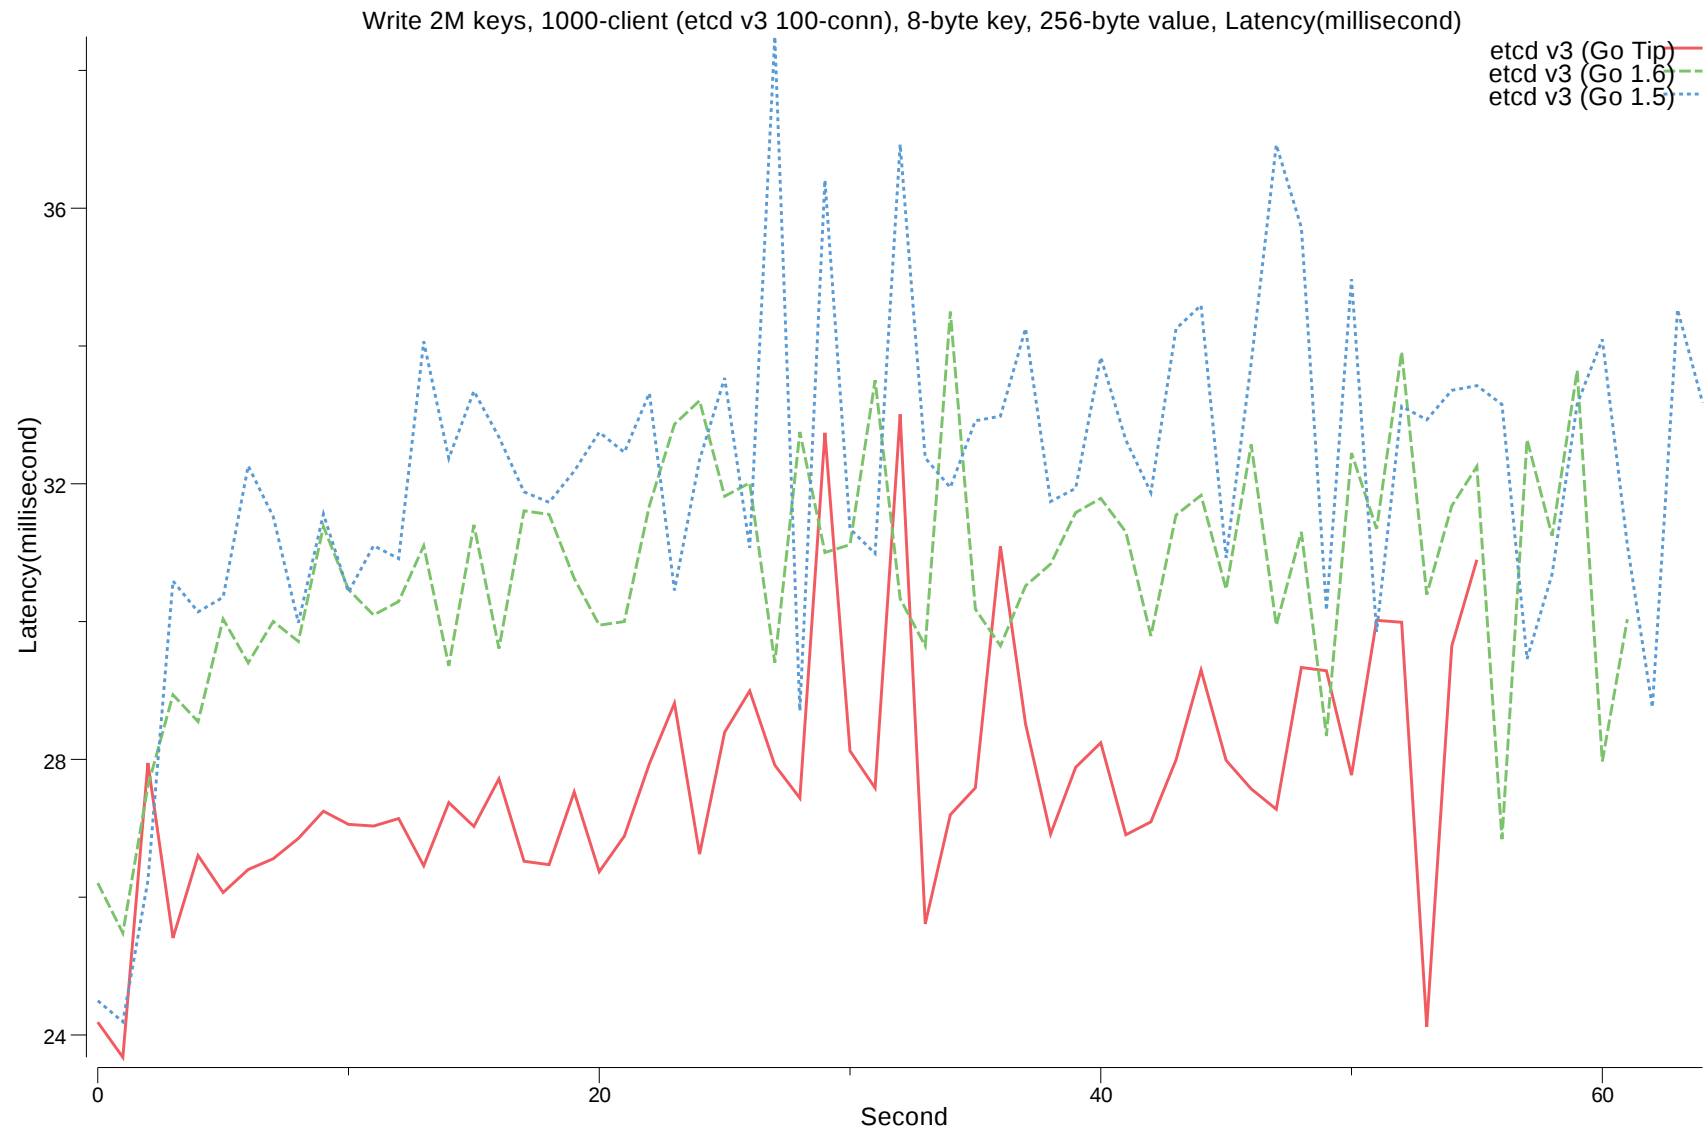
etcd Go 10 Tips

#1 Use latest Go

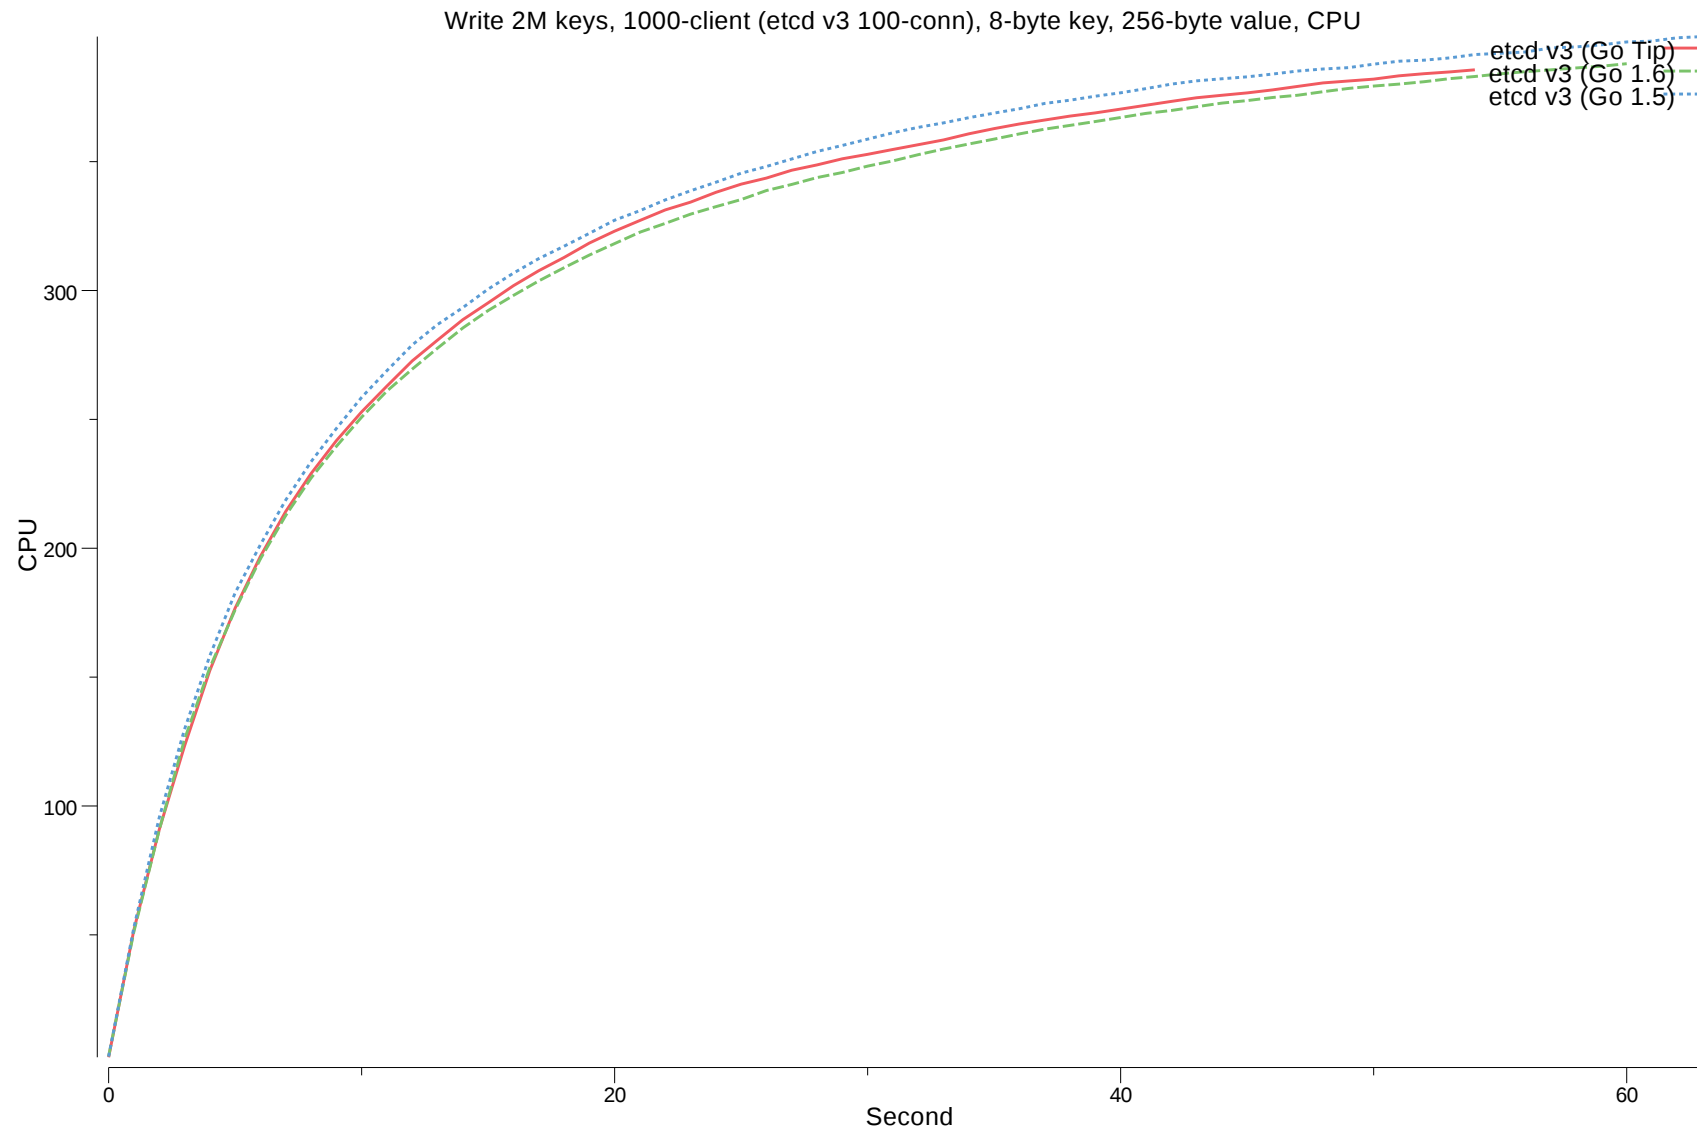
Throughput



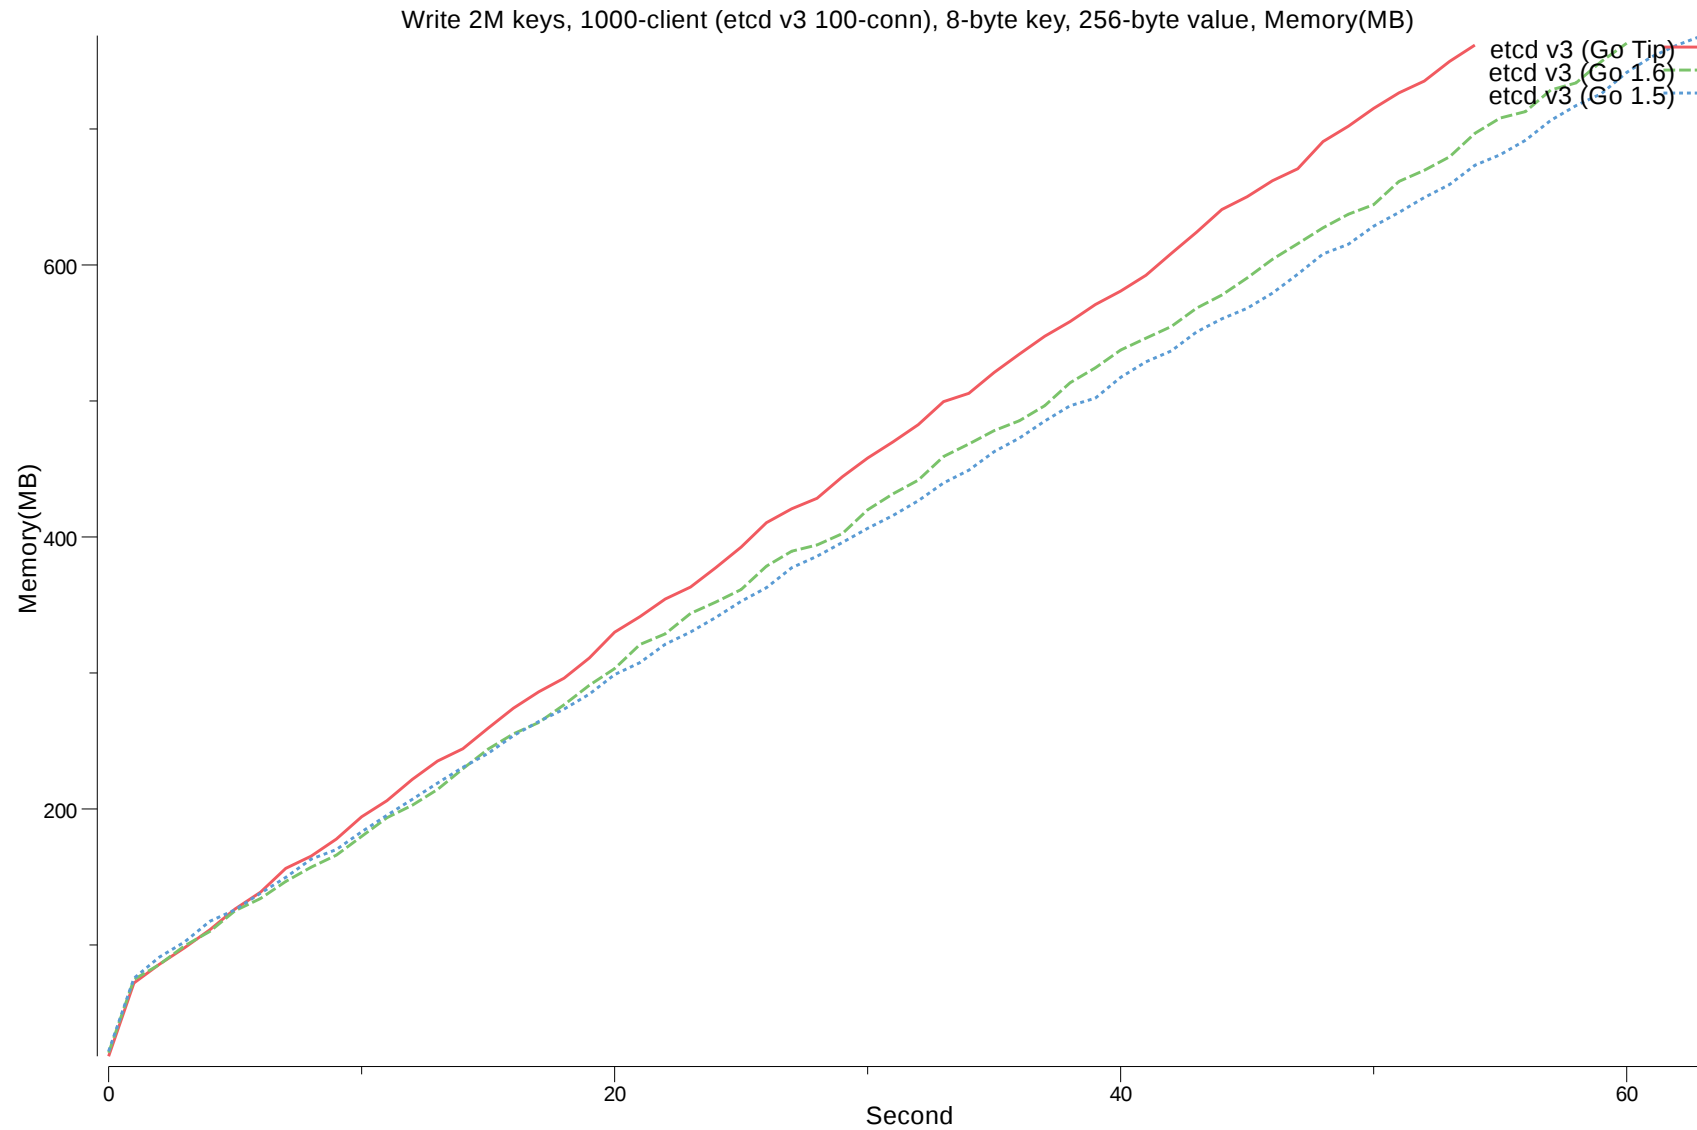
Latency



CPU



Memory



#2 Check slice allocation

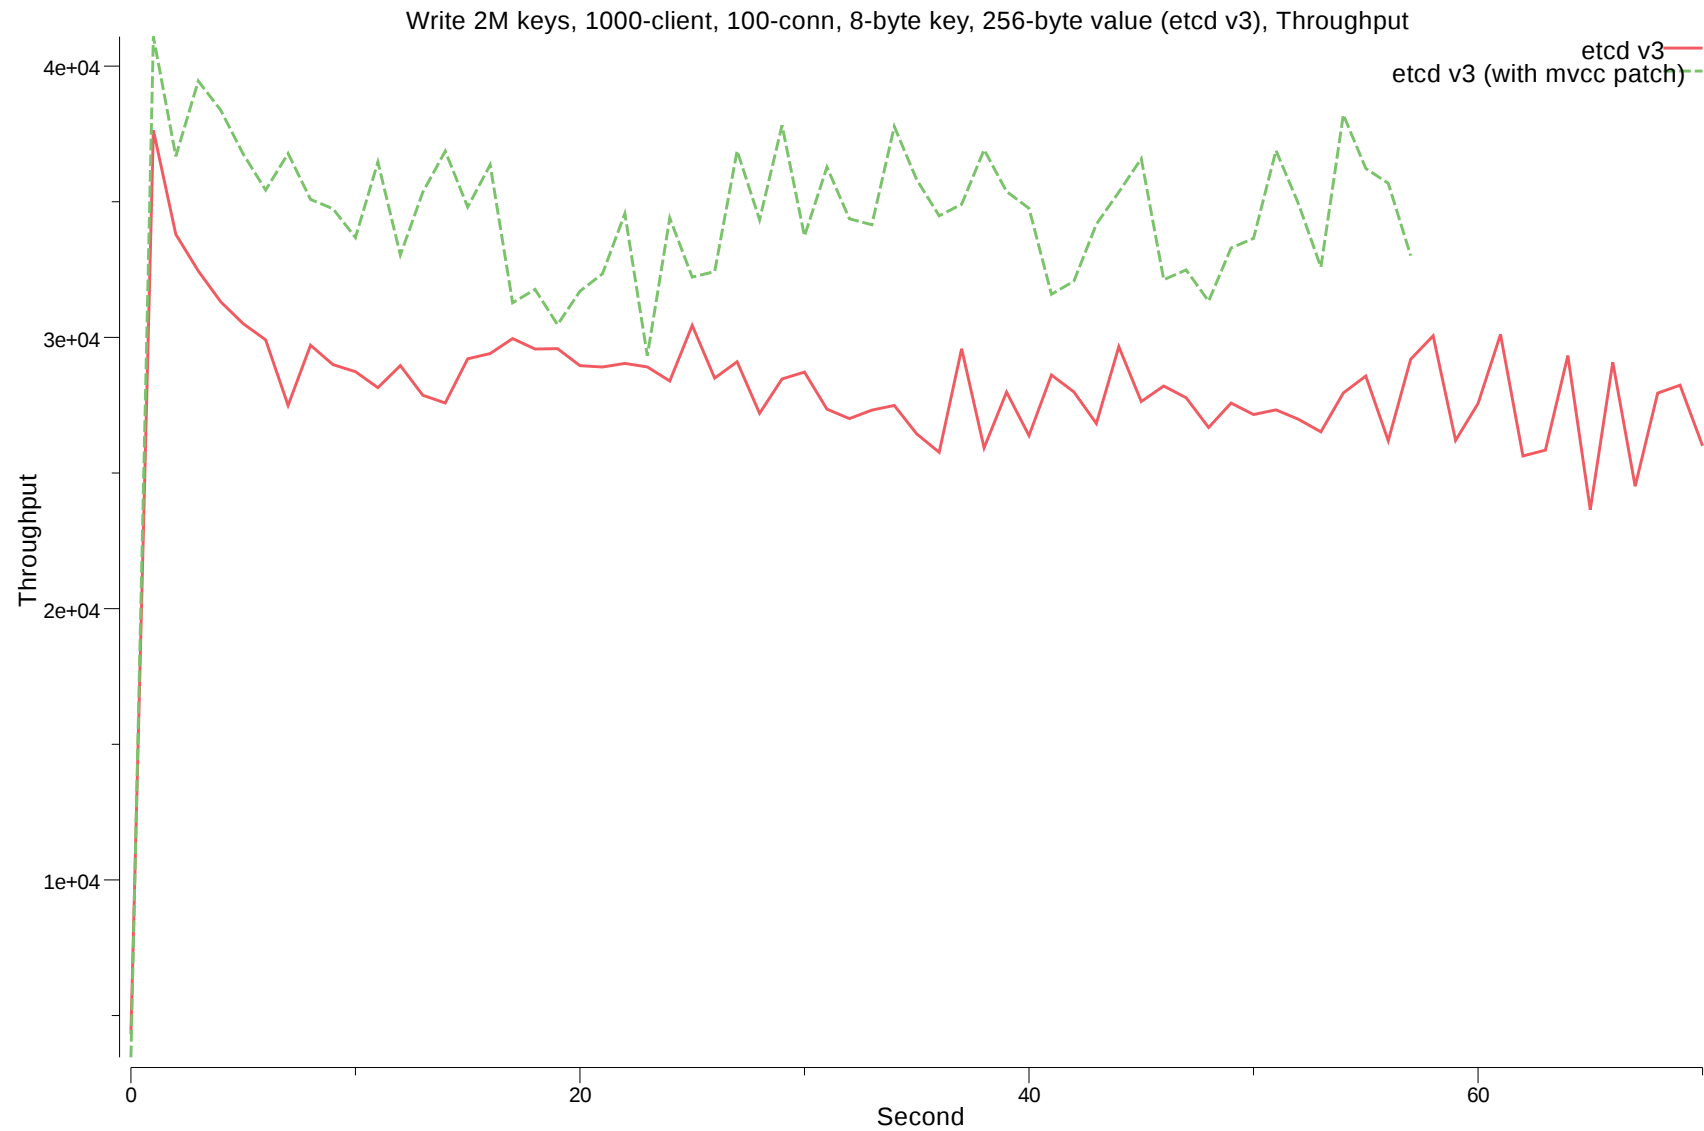
[GH5238](https://github.com/coreos/etcd/pull/5238) (<https://github.com/coreos/etcd/pull/5238>)

```
s.changes = make([]mvccpb.KeyValue, 0, 128)
s.changes = make([]mvccpb.KeyValue, 0, 4)    // better for etcd use case
```

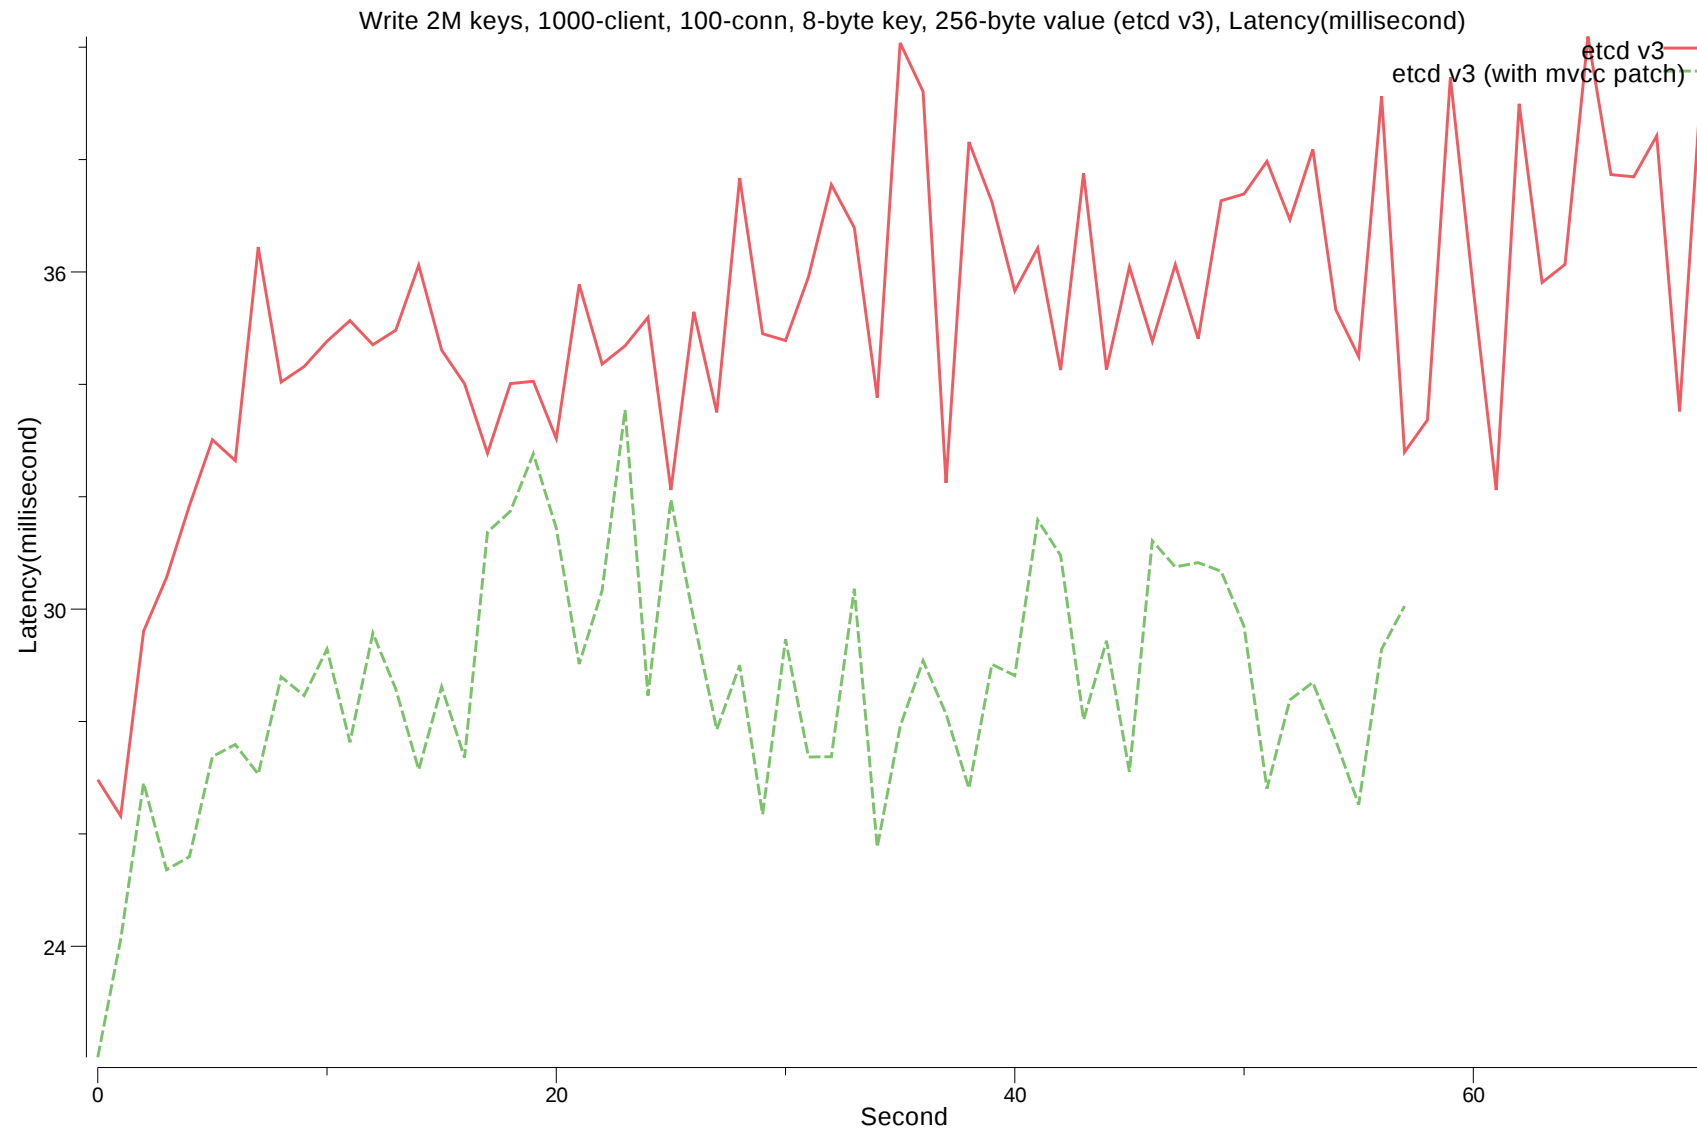
More capacity means more allocation.

Check your use case when making slice!

Throughput



Latency



#3 Test your code

All with "go test" command

- unit tests
- integration tests
- functional tests
- benchmarks

Use expect package for e2e tests

```
proc, _ = expect.NewExpect("etcdctl", "get", "foo")  
_, err = proc.Expect("bar") // if err != nil, found a bug!
```

github.com/coreos/etcd/pkg/expect (<https://godoc.org/github.com/coreos/etcd/pkg/expect>)

#4 Check goroutine leaks

Scan runtime.Stack

```
func TestMain(m *testing.M) {  
    v := m.Run()  
    if v == 0 && testutil.CheckLeakedGoroutine() {  
        os.Exit(1)  
    }  
    os.Exit(v)  
}  
  
func TestSample(t *testing.T) {  
    defer testutil.AfterTest(t)  
    ...  
}
```

- [net/http/main_test.go](https://github.com/golang/go/blob/master/src/net/http/main_test.go) (https://github.com/golang/go/blob/master/src/net/http/main_test.go)
- [github.com/coreos/etcd/pkg/testutil](https://godoc.org/github.com/coreos/etcd/pkg/testutil) (<https://godoc.org/github.com/coreos/etcd/pkg/testutil>)

Highly recommend for projects with context.Context, gRPC

#5 Always gofmt, go vet

gofmt

```
Checking gofmt...
gofmt checking failed:
version/a.go
diff version/a.go gofmt/version/a.go
--- /tmp/gofmt6613415602016-05-15 04:07:11.087869561 +0000
+++ /tmp/gofmt2762292392016-05-15 04:07:11.087869561 +0000
@@ -15,5 +15,6 @@
 package version

 func myFunc() {
- a := 1
- a += 1 }
+     a := 1
+     a += 1
+ }
```

go vet

```
log.Fatalf("hello %d", "a")
// arg "a" for printf verb %d of wrong type: string
```


#6 Write simple Go

```
ok := true
if ok == true {} // X
if ok {}         // 0
```

Don't:

```
err := l.newStream()
if err != nil {
    return err
}
return nil
```

Do:

```
return l.newStream()
```

github.com/dominikh/go-simple by Dominik

#7 Check unused

- github.com/dominikh/go-unused (<https://github.com/dominikh/go-unused>) by Dominik
- Finds unused constants, variables, functions and types

```
func reportMetrics() {}  
// func reportMetrics is unused
```

Found bugs in etcd [GH4955](https://github.com/coreos/etcd/pull/4995/files) (<https://github.com/coreos/etcd/pull/4995/files>)

#8 Use goword

- github.com/chzchzchz/goword (https://github.com/chzchzchz/goword) by Anthony (etcd team)

Comment checker

```
// This.  
func Hello() {} // This. (godoc-export: This -> Hello?)"
```

Spell checker

```
// Hello retuens.  
func Hello() {} // Hello retuens. (spell: retuens -> returns?)
```

#9 Document with godoc

- etcd must be easy to use
- etcd needs good documentation
- Example? [etcd/clientv3](https://github.com/coreos/etcd/tree/master/clientv3) (<https://github.com/coreos/etcd/tree/master/clientv3>)

#10 vendor

Problem

- etcd client package is used within etcd repo (etcdctl)
- etcd client imports gRPC and vendors it
- Project B import this etcd client package
- Project B also uses gRPC but from different import path

Now two projects has conflicting gRPC code [GH566](https://github.com/grpc/grpc-go/issues/566) (<https://github.com/grpc/grpc-go/issues/566>)

```
panic: http: multiple registrations for /debug/requests
```

#10 vendor

Solution [GH4950](https://github.com/coreos/etcd/pull/4950) (<https://github.com/coreos/etcd/pull/4950>)

- Create symlinks inside cmd directory
- In -s main.go cmd/main.go
- cmd/vendor

Update dependency

- In -s cmd/vendor vendor
- godep save

Still go-get-able. No conflicts with other projects.

Works, even on Windows!

Thank you

Gyu-Ho Lee

CoreOS

gyu_ho.lee@coreos.com (mailto:gyu_ho.lee@coreos.com)

<https://github.com/coreos/etcd> (https://github.com/coreos/etcd)

